

# Principles of designing Go APIs with channels



Alan Shreve  
Keen IO  
[@inconsreveable](#)

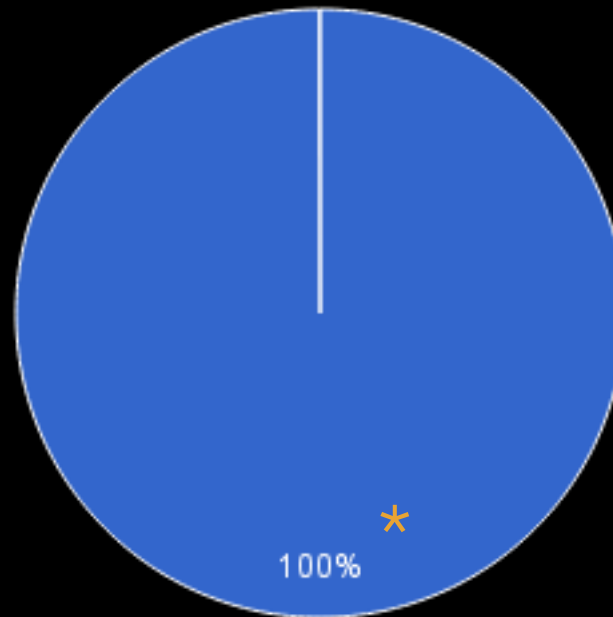
# First-class concurrency support in Go



- APIs in stdlib without `chan`: ~ 30,000
- APIs in stdlib with `chan`: ~10

# APIs in the Standard Library

Sync vs Async APIs



\*rounded up

# io.Reader

```
type Reader {  
    Read([]byte) (int, error)  
}
```

# Where is `io.AsyncReader`?

```
type ReadResult struct {  
    n      int  
    err    error  
}  
  
type AsyncReader {  
    ReadAsync([]byte) <-chan ReadResult  
}
```

APIs should be synchronous (blocking).


*Principle 0*



- Concurrency in Go is easy
- Leave concurrency decisions to the caller
  - Easier for you
  - More flexible for the caller
- Reduces API surface area (no additional Async APIs)
- *Your API still needs to be thread-safe.*
- **Most important takeaway of this talk**



```
n, err = conn.Read(buf)
```



```
reads := make(chan readResult)
go func() {
    n, err = conn.Read(buf)
    reads <- readResult{n, err}
}()
```

So - when is it  
appropriate to have a  
channel in your API?

let's come back to this later

An API should declare the directionality of its channels.

*Principle 1*



```
func After(d Duration) <-chan Time
```

```
func Notify(c chan<- os.Signal,  
           sig ...os.Signal)
```

“The optional <- operator specifies the channel direction, send or receive. If no direction is given, the channel is bidirectional.”

– *The Go Programming Language Specification*


- Proper API usage enforced by the compiler
- Type signature elucidates API data flow and proper usage

# This will not compile

```
t := time.After(time.Second)
```

```
t <- time.Now()
```

```
send to receive-only type <-chan Time
```



An API that sends an **unbounded** stream of values into a channel **must** document how it behaves for slow consumers.

*Principle 2*



```
// NewTicker returns a new Ticker containing
// a channel that will send the
// time with a period specified by the
// duration argument.
// It adjusts the intervals or drops ticks
// to make up for slow receivers.
// ...
func NewTicker(d Duration) *Ticker {
    ...
}
```

```
// Notify causes package signal to relay
// incoming signals to c.
// ...
// Package signal will not block sending to c
// ...
func Notify(c chan<- os.Signal, sig ...os.Signal) {
    ...
}
```

```
// OpenChannel tries to open an channel.  
// ...  
// On success it returns the SSH Channel  
// and a Go channel for incoming, out-of-band  
// requests. The Go channel must be serviced, or  
// the connection will hang.  
OpenChannel(name string, data []byte) (Channel, <-  
chan *Request, error)
```

- Implementation has a choice when sending values into a channel that is full:
  - Block
  - Don't block

```
select {  
  case c<-val:  
  default:  
}
```
- No language annotation, must be documented.



An API that sends a **bounded** set of values into a channel it *accepted as an argument* **must** document how it behaves for slow consumers.

*Principle 3*

```
func (client *Client) Go(serviceMethod string,  
                           args interface{},  
                           reply interface{},  
                           done chan *Call  
                           ) *Call
```

- Only sends **one** value into the channel
- Channel was an argument, could still be full
- No documented behavior

```
171 func (call *Call) done() {
172     select {
173     case call.Done <- call:
174         // ok
175     default:
176         // We don't want to block here.
177         // It is the caller's responsibility to make
178         // sure the channel has enough buffer space.
179         // See comment in Go().
180         if debugLog {
181             log.Println("rpc: discarding Call reply due
to insufficient Done chan capacity")
182         }
183     }
184 }
```





An API that sends a **bounded** set of values may do so safely by returning an appropriately buffered channel.

*Principle 4*

```
type CloseNotifier interface {  
    // CloseNotify returns a channel that  
    // receives a single value  
    // when the client connection has gone  
    // away.  
    CloseNotify() <-chan bool  
}
```

```
func After(d Duration) <-chan Time
```

- Returned channel can be buffered to maximum number of values to send
- Sending on the channel guaranteed to never block

# Act II

# Tradeoffs

An API sending an **unbounded** stream of values must trade off between *accepting the channel as an argument* and *returning a new channel*.

*Tradeoff 0*

```
func Notify(c chan<- os.Signal, sig ...os.Signal)
```

**VS**

```
func Notify(sig ...os.Signal) <-chan os.Signal
```

- why doesn't `signal.Notify` allocate the channel for you?
- would have to choose channel buffer size
  - memory tradeoff
  - tradeoff tolerance of missed signals
- why not this then:
  - ```
func Notify(sig ...os.Signal, size int) <-chan os.Signal
```
  - caller loses ability to use the same channel for handling (one goroutine)



```
sigs := make(chan os.Signal, 10)
go handleSigs(sigs)
signal.Notify(sigs, os.Interrupt)

// later ...
signal.Notify(sigs, os.Kill)
```

but on the other  
hand . . .

```
func NewClientConn(c net.Conn,  
                  addr string,  
                  config *ClientConfig)  
(Conn, <-chan NewChannel, <-chan *Request, error)
```

**VS**

```
func NewClientConn(c net.Conn,  
                  addr string,  
                  config *ClientConfig,  
                  newChans chan<- NewChannel,  
                  newReqs chan<- *Request)  
(Conn, error)
```

- Returned channel makes it easy to know when an SSH connection is finished
- The channels close - need another mechanism when channel is an argument
- Returned channel grants compiler protection against send on closed channel

```
var c <-chan  
close(c)
```

- `close(c)` (cannot close receive-only channel)

So - when is it  
appropriate to have a  
channel in your API?

What do these APIs all  
have in common?

```
func After(d Duration) <-chan Time
```

```
func Notify(c chan<- os.Signal,  
           sig ...os.Signal)
```

```
type CloseNotifier interface {  
    CloseNotify() <-chan bool  
}
```

```
func (client *Client) Go(serviceMethod string,  
                        args interface{},  
                        reply interface{},  
                        done chan *Call  
                        ) *Call
```

- When a timer fires
- When the process receives a signal
- When a client closes the HTTP connection
- When a response is received to an RPC call
- When a new SSH channel is opened



Asynchronous  
notification of events

# A challenger appears

```
func ListenAndServe(addr string,  
                    handler Handler) error
```

# Callbacks, in my Go code?

It's more likely than you think.

Any API written with a channel could use a callback instead.

*Tradeoff 1*

# Mirror Universe



# Mirror Universe

```
func ListenAndServe(addr string)
    (chan *http.Resp, error)
```

```
func Notify(handler func(os.Signal),
    sig ...os.Signal)
```

# Tradeoffs

- Callback overhead is a function call; channel has synchronization overhead
- Callbacks can be wrapped (recover from panic, behavior after completion)
- Trivial to make a callback that sends into a channel
- Interface callback make for easy composition

# Tradeoffs

- Channels are more idiomatic than callbacks
- Channels make code-flow easier to reason about
- I just left Node.js . . .



# Takeaways

- Always prefer synchronous APIs if possible
- Declare the directionality of your channels
- Document the behavior of your APIs in the presence of slow consumers

# Takeaways

- Weigh tradeoff between channel as an argument vs channel as a return value
- Consider using callbacks instead of channels for notifying your caller of async events



# Questions

Alan Shreve  
Keen IO  
[@inconshreveable](#)