


High performance git infrastructure with Go


David Calavera
Code Climate



git

2005






2008



2010...



Travis CI




How do we make git faster?

```
system("git log ...")*
```



```
git fetch origin
```



libgit2

Disclaimer: examples do not handle errors 😱

github.com/libgit2/git2go

```
// Load with side effects.  
// Initialize libgit2's TLS:  
//  
// func init() {  
//     C.git_libgit2_init()  
// }  
//  
// Import package “git”, which is  
// not very goimports friendly.  
import “github.com/libgit2/git2go”
```

```
// Create a new repository.  
// Do not use a working directory.  
path := "/var/git/repository"  
bare := true  
r, _ := git.InitRepository(path, bare)  
  
// Clone a repository.  
url := "git://github.com/golang/go"  
opts := git.CloneOptions{Bare: bare}  
r, _ := git.Clone(url, path, &opts)
```

```
// Create a new remote ref.  
name    := "my-fork"  
url     := "git://github.com/wadus/go"  
rm, _   := r.CreateRemote(name, url)
```

```
// Fetch all refs from a remote.  
var refsspecs []string  
rm.Fetch(refspecs, nil, nil)
```

```
// Search for objects.
```

```
sha := "4c279186e24f7b3a59aa682a870747df6eaca013"
```

```
oid := git.NewOid(sha)
```

```
c, _ := r.LookupCommit(oid)
```

```
b, _ := r.LookupBlob(oid)
```

```
t, _ := r.LookupTree(oid)
```

```
o, _ := r.Lookup(oid)
```

```
fmt.Printf("👩 %v\n", o.Type())
```

```
// Read commit data.  
sha := "4c279186e24f7b3a59aa682a870747df6eaca013"  
oid := git.NewOid sha  
path := "src/os/exec.go"  
  
c, _ := r.LookupCommit(oid)  
t, _ := c.Tree()  
e, _ := t.EntryByPath(path)  
  
b, _ := r.LookupBlob(e.Id())  
fmt.Printf("🎉 %q\n", b.Contents())
```




```
// Commit new changes.  
idx, _ := r.Index()  
idx.AddByPath("src/os/exec.go")  
t, _ := idx.WriteTree()  
idx.Write()  
  
h, _ := r.Head()  
c, _ := r.LookupCommit(h)  
  
s := &git.Signature{"me", "me@me.com", time.Now()}  
m := "Add moar changes"  
r.CreateCommit("", s, s, m, t, c)
```

The average latency was also lower with [REDACTED]:

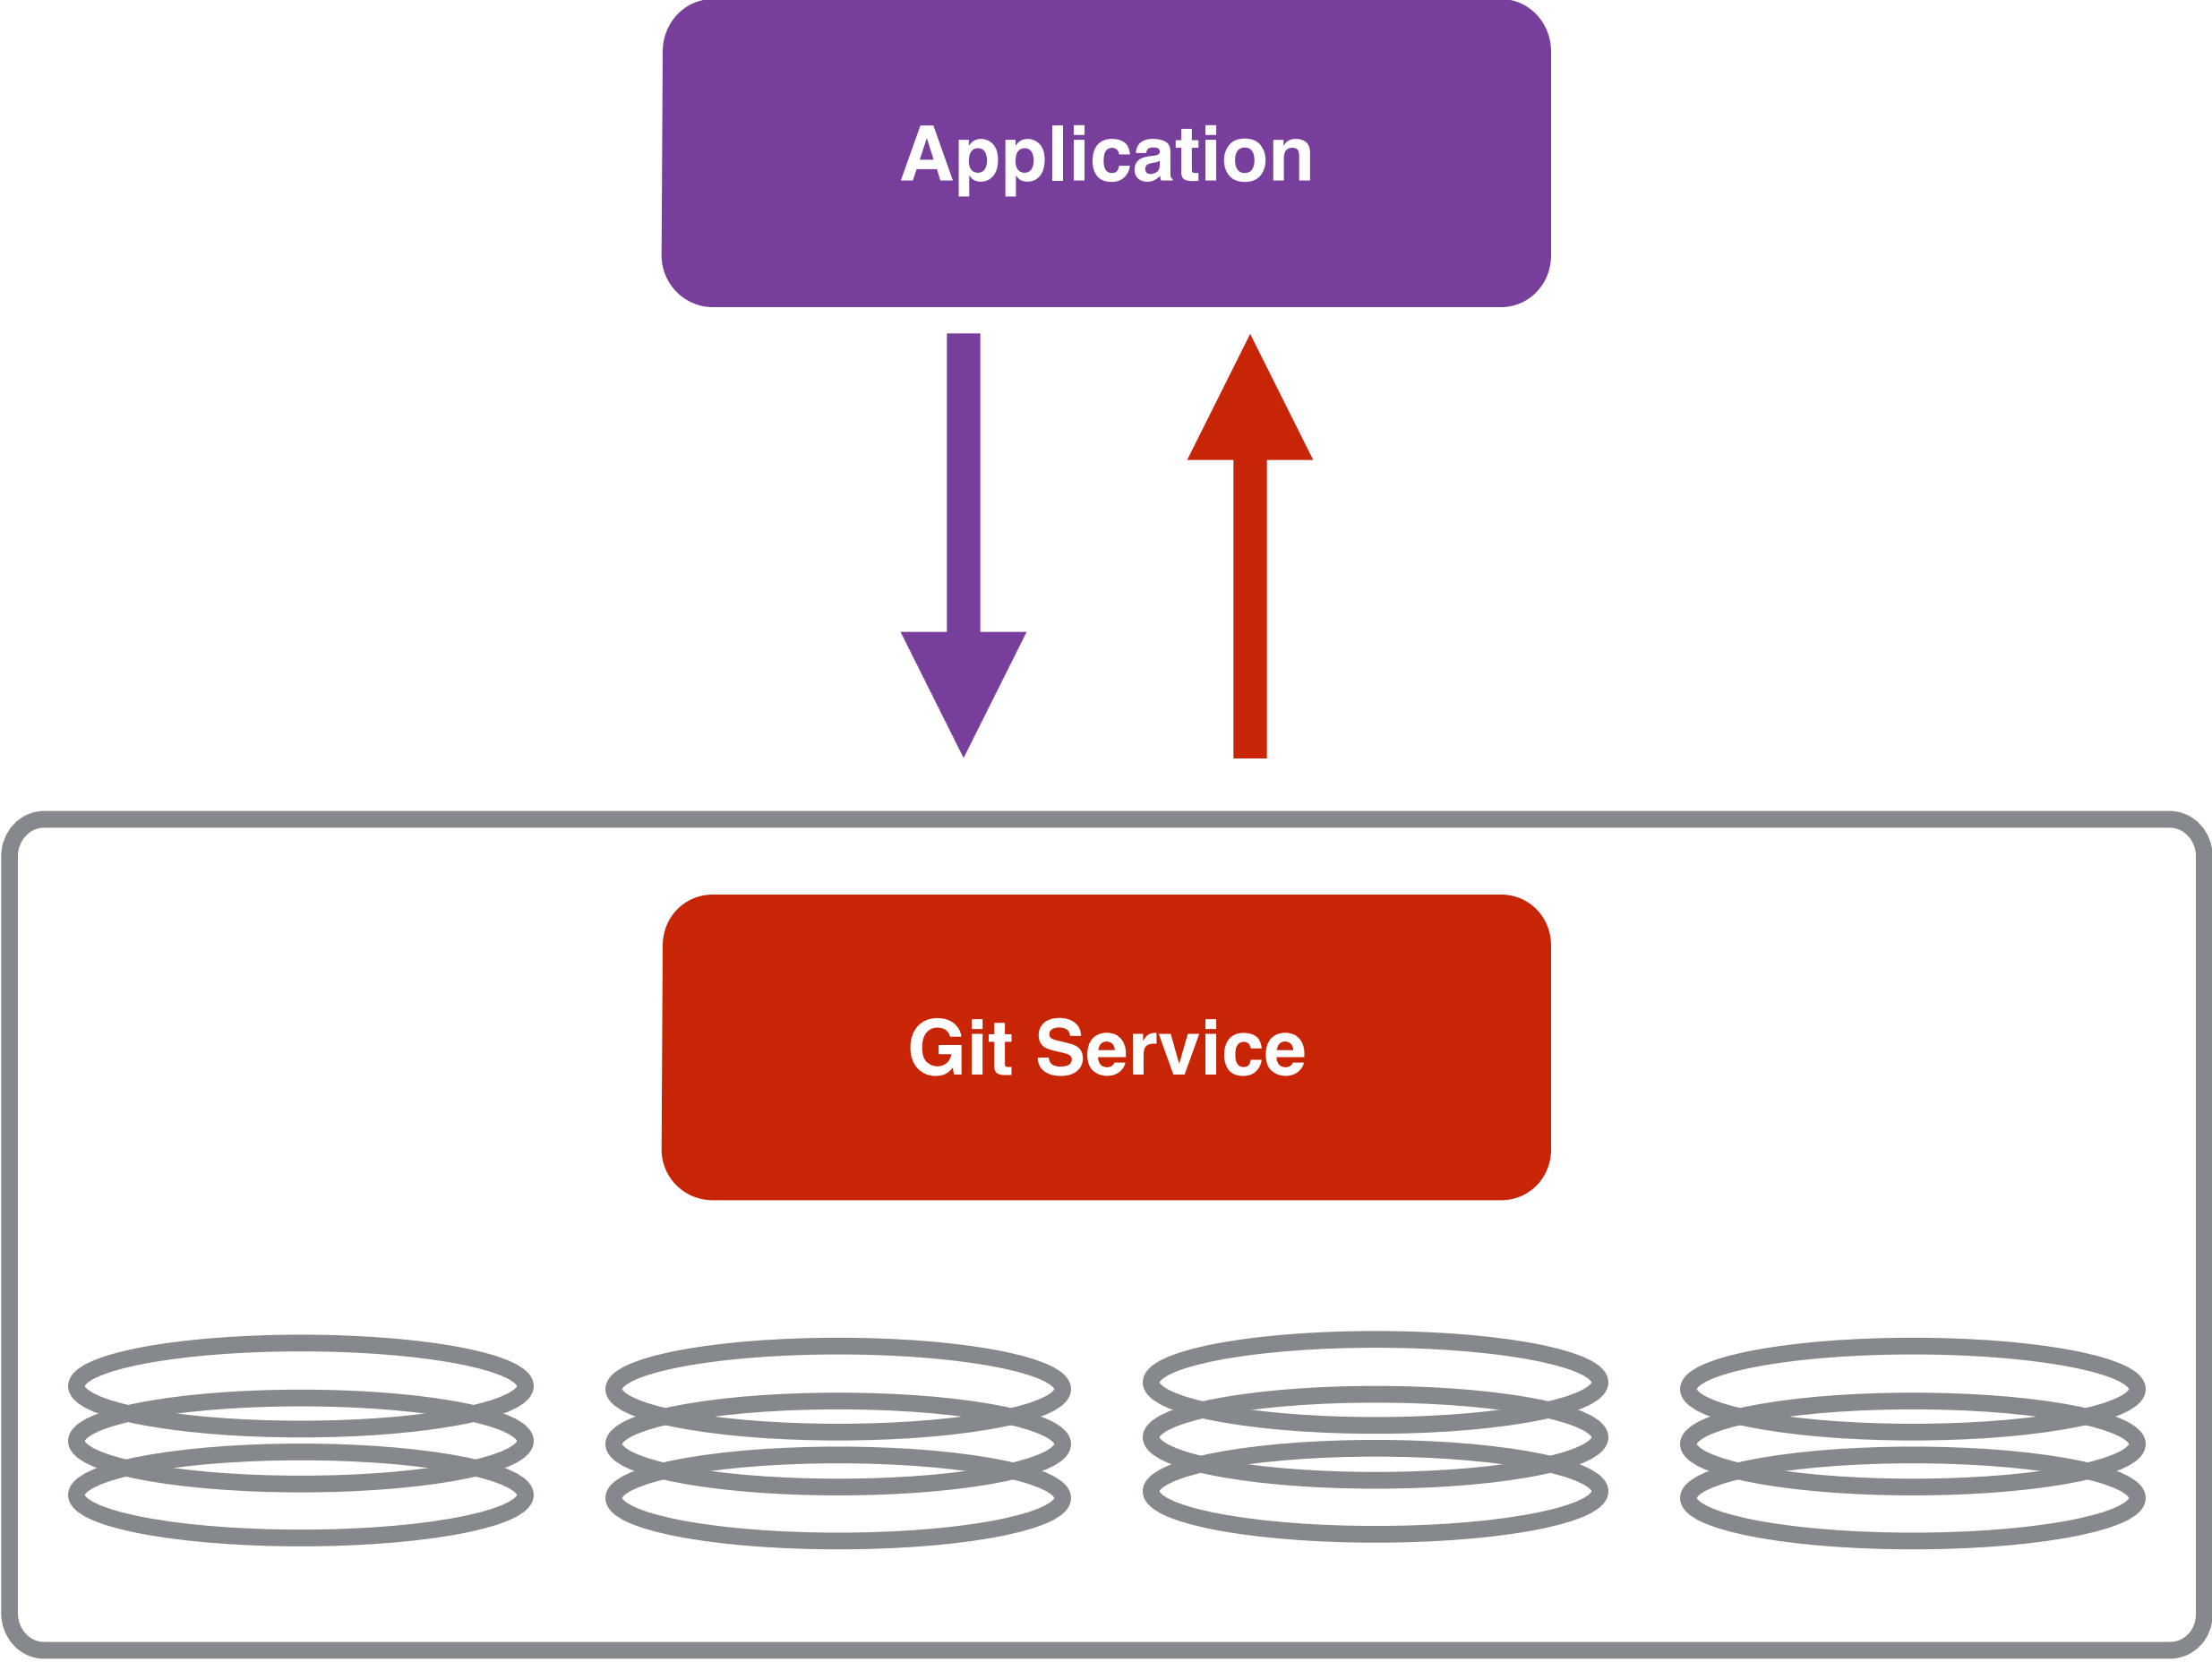


I'm removing the experiment and merging these changes.

Memory in use



Designing a distributed git storage



Constraint your data model

```
// protocol buffers schema.
```

```
message Branch {  
    required string name = 1;  
}
```

```
message Repository {  
    optional string name = 1;  
    repeated Branch branches = 2;  
}
```



```
// Read branches.
var branches []*pb.Branch

f := func(b *git.Branch, t git.BranchType) error {
    n, _ := b.Name()
    p := &pb.Branch{
        Name: &n,
    }
    branches = append(branches, p)
    return nil
}

b, _ := r.NewBranchIterator(git.BranchRemote)
b.ForEach(f)
```



```
// Read branches via http.  
h := func(w http.ResponseWriter, r *http.Request) {  
    pbBranches := readBranches(r)  
    pbRepo      := &pb.Repository{  
        Branches: pbBranches,  
    }  
  
    data, _ := proto.Marshal(pbRepo)  
    w.Write(data)  
}  
  
http.HandleFunc("/r/foo/branches", h)
```

Design from first principles

*A shared-data system can have at most two of the three following properties: **C**onsistency, **A**vailability, and tolerance to network **P**artitions*

Dr. Eric Brewer

*You Can't Sacrifice
Partition Tolerance*

Coda Hale

github.com/afex/hystrix-go

github.com/rubyist/circuitbreaker

github.com/eapache/go-resiliency/braker

```
// Read branches via http.
import "github.com/rubyist/circuitbreaker"

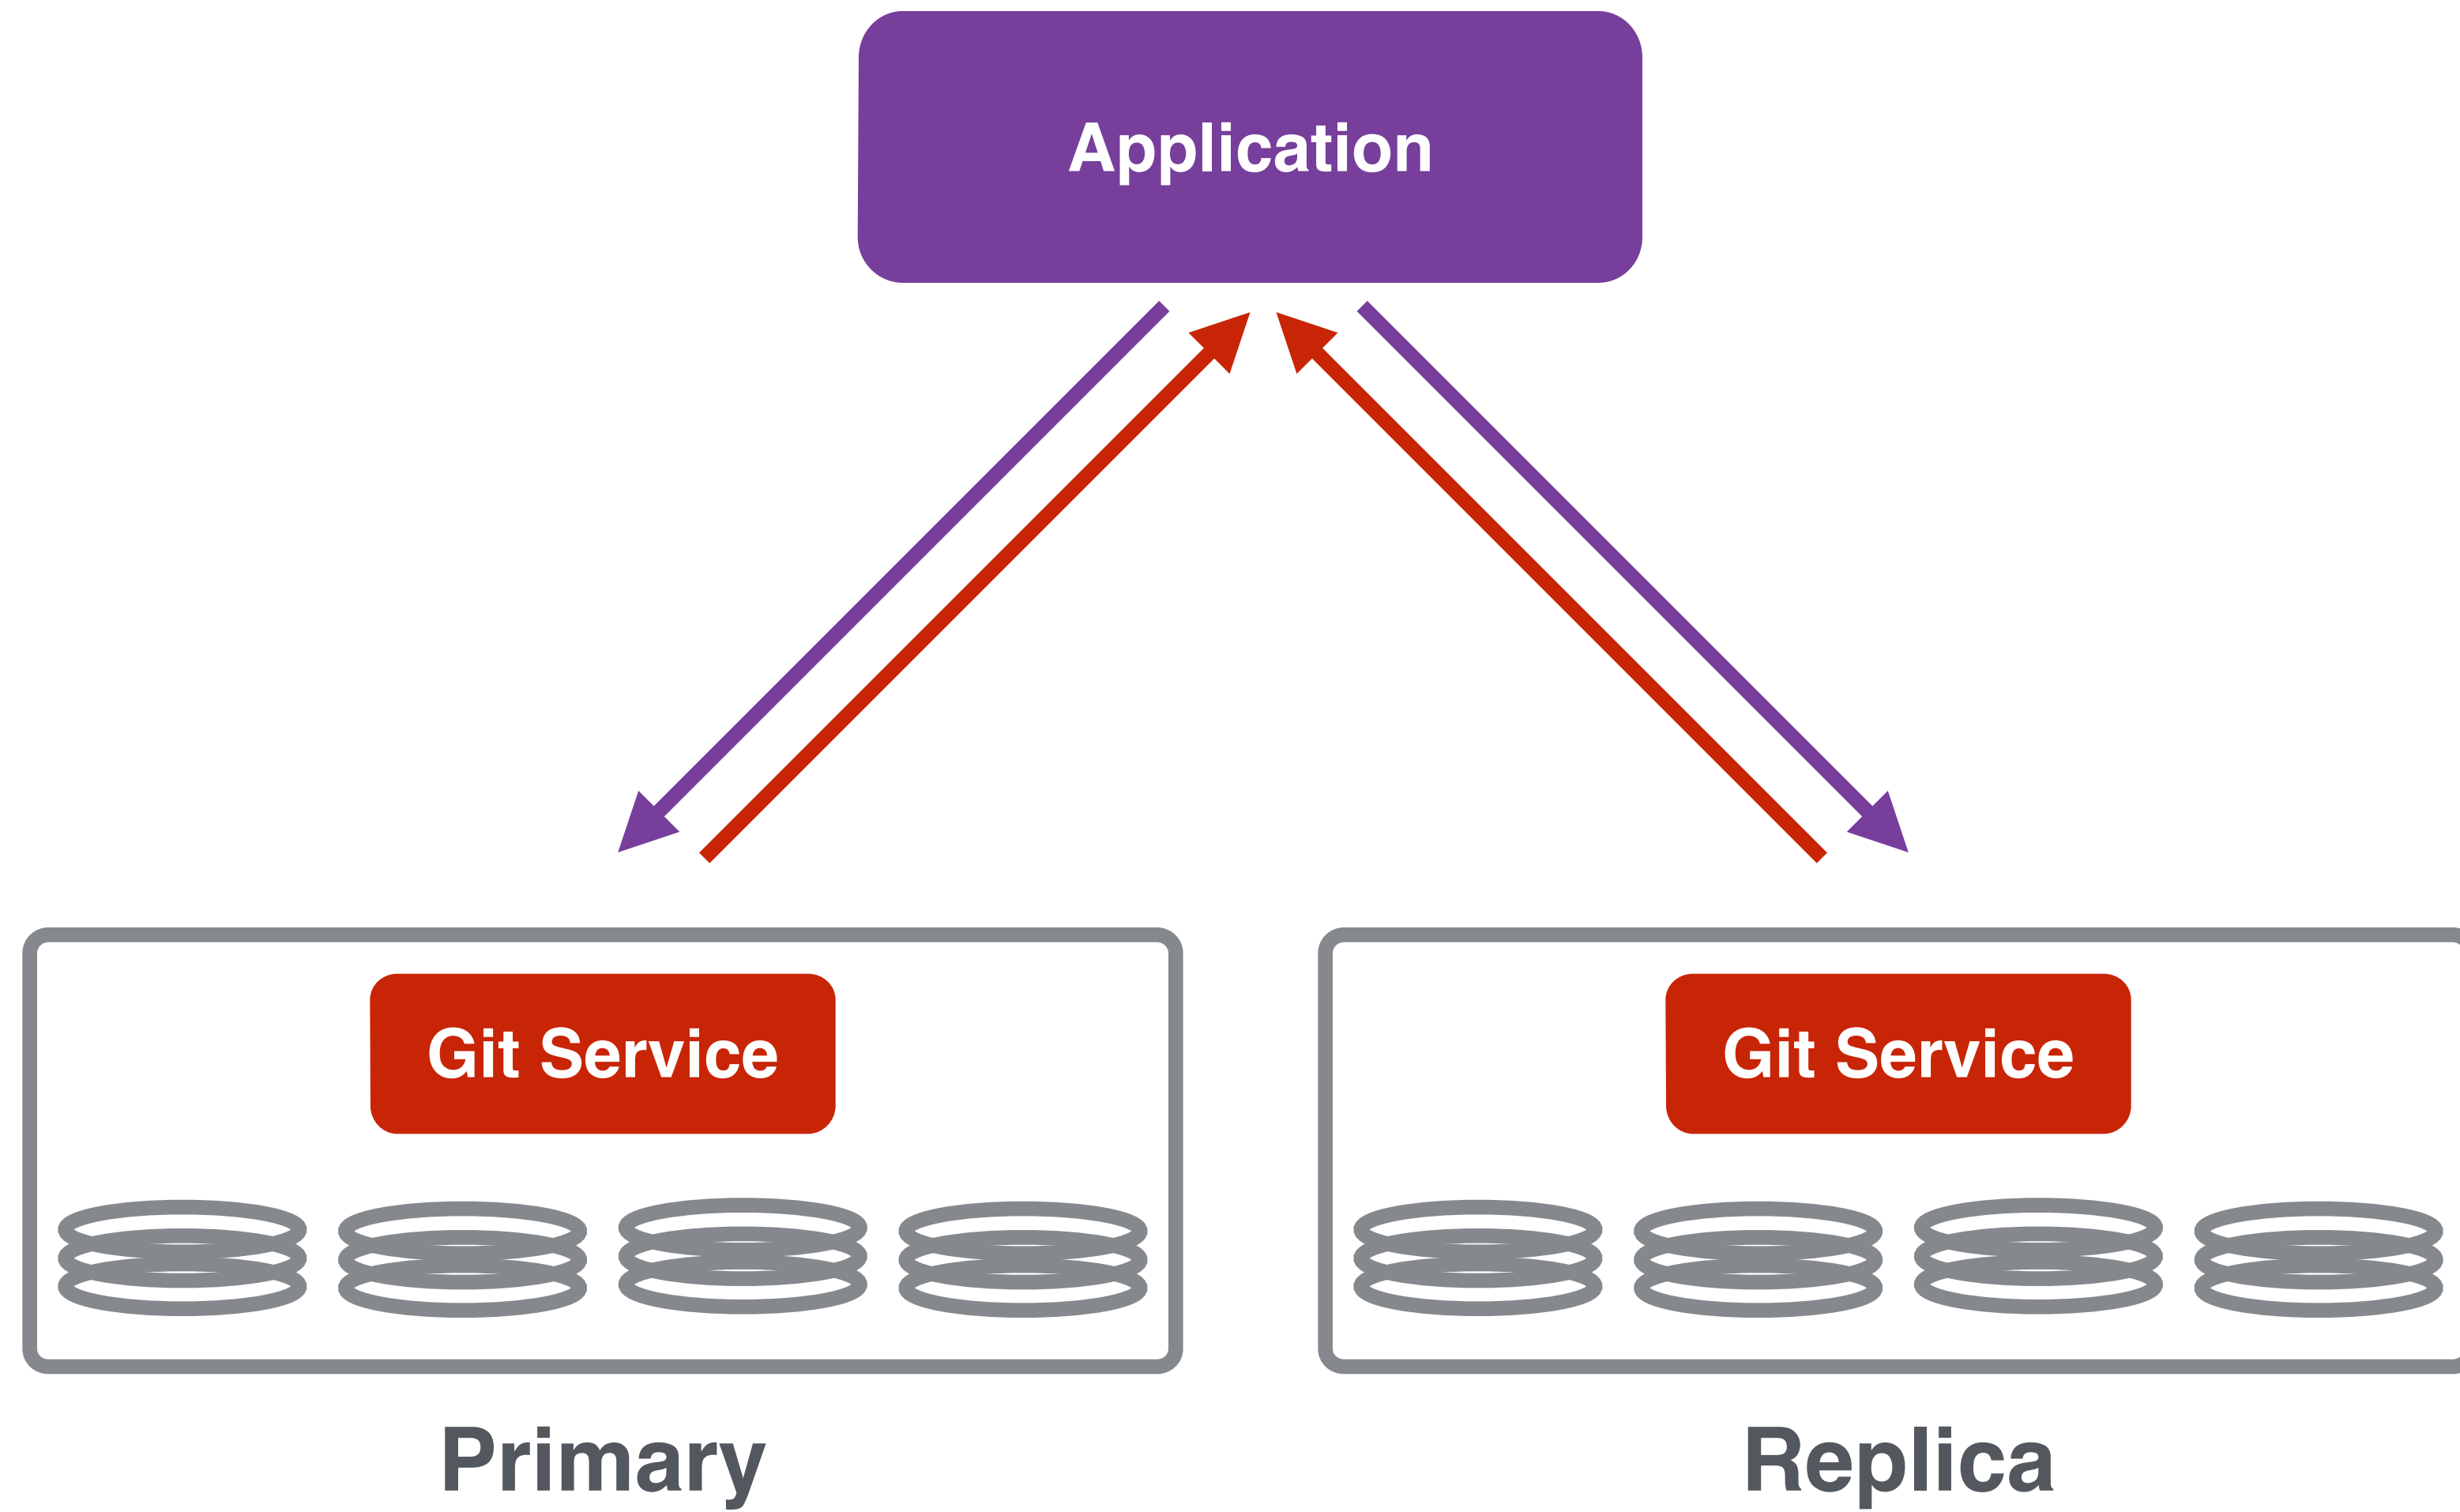
out := 5 * time.Second
url := "http://git-server/r/foo/branches"
c    := circuit.NewHTTPClient(out, 10, nil)

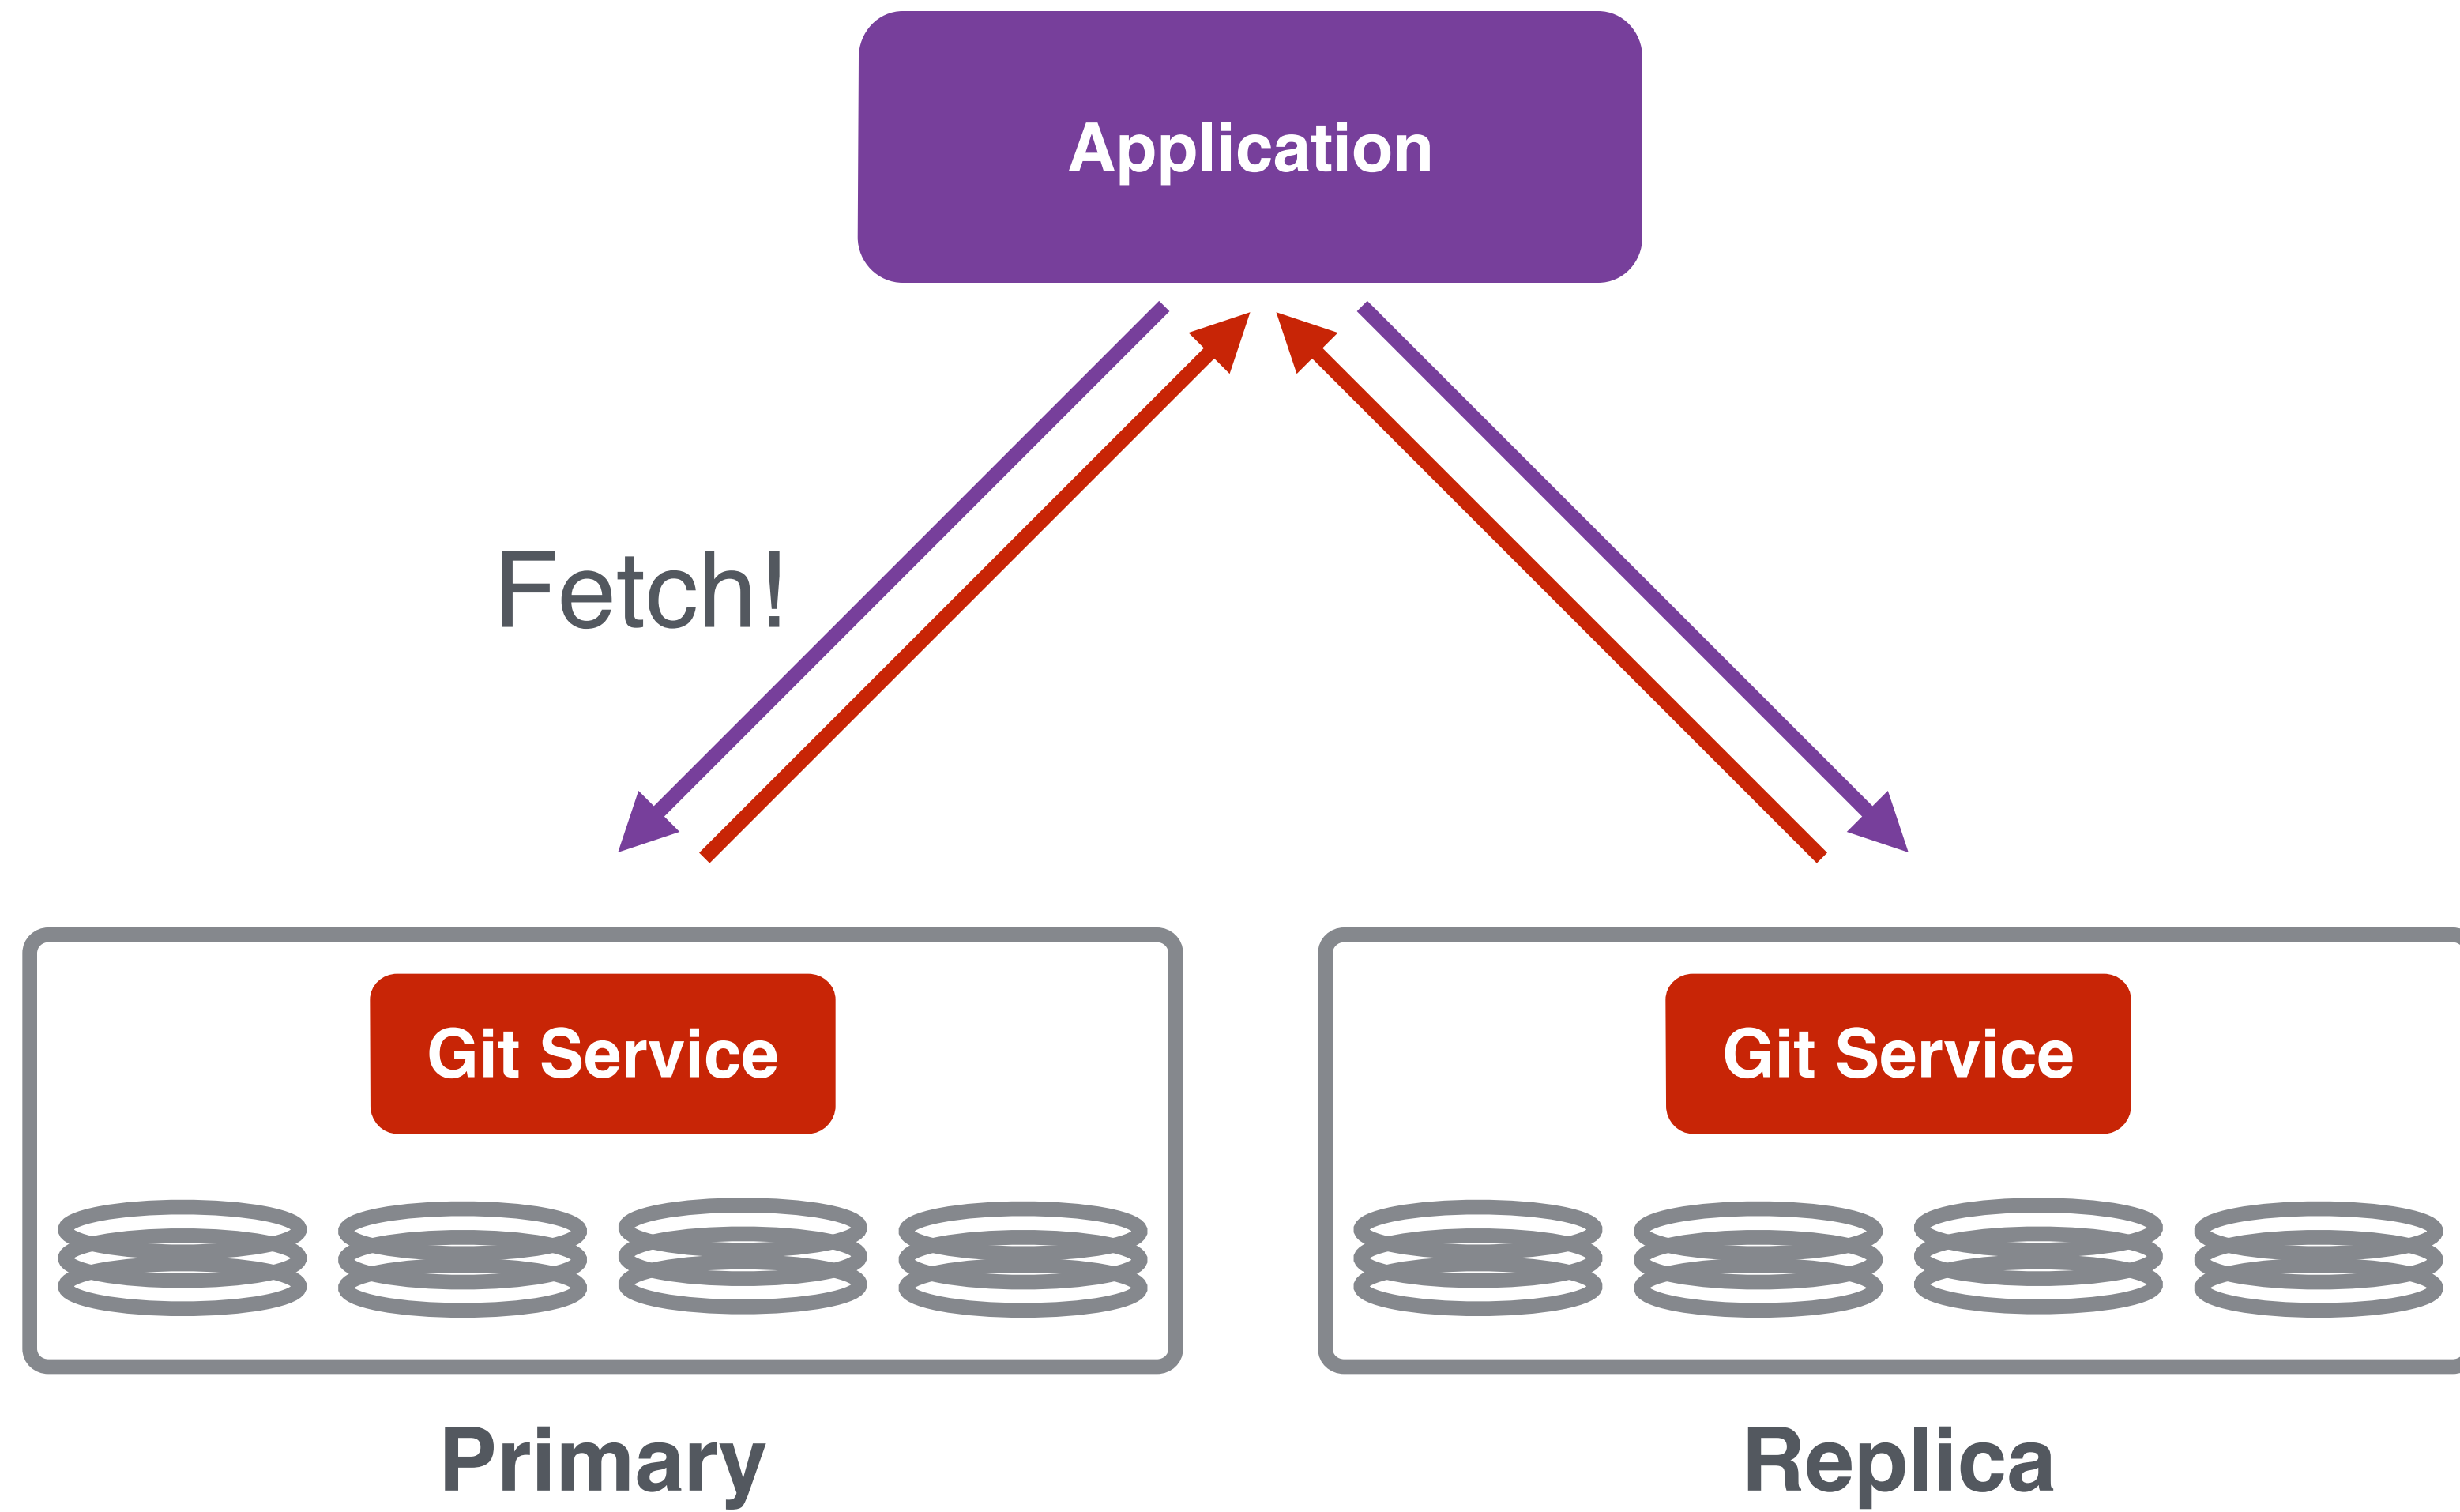
c.BreakerTripped = func() {
    // Handle partition error response.
}

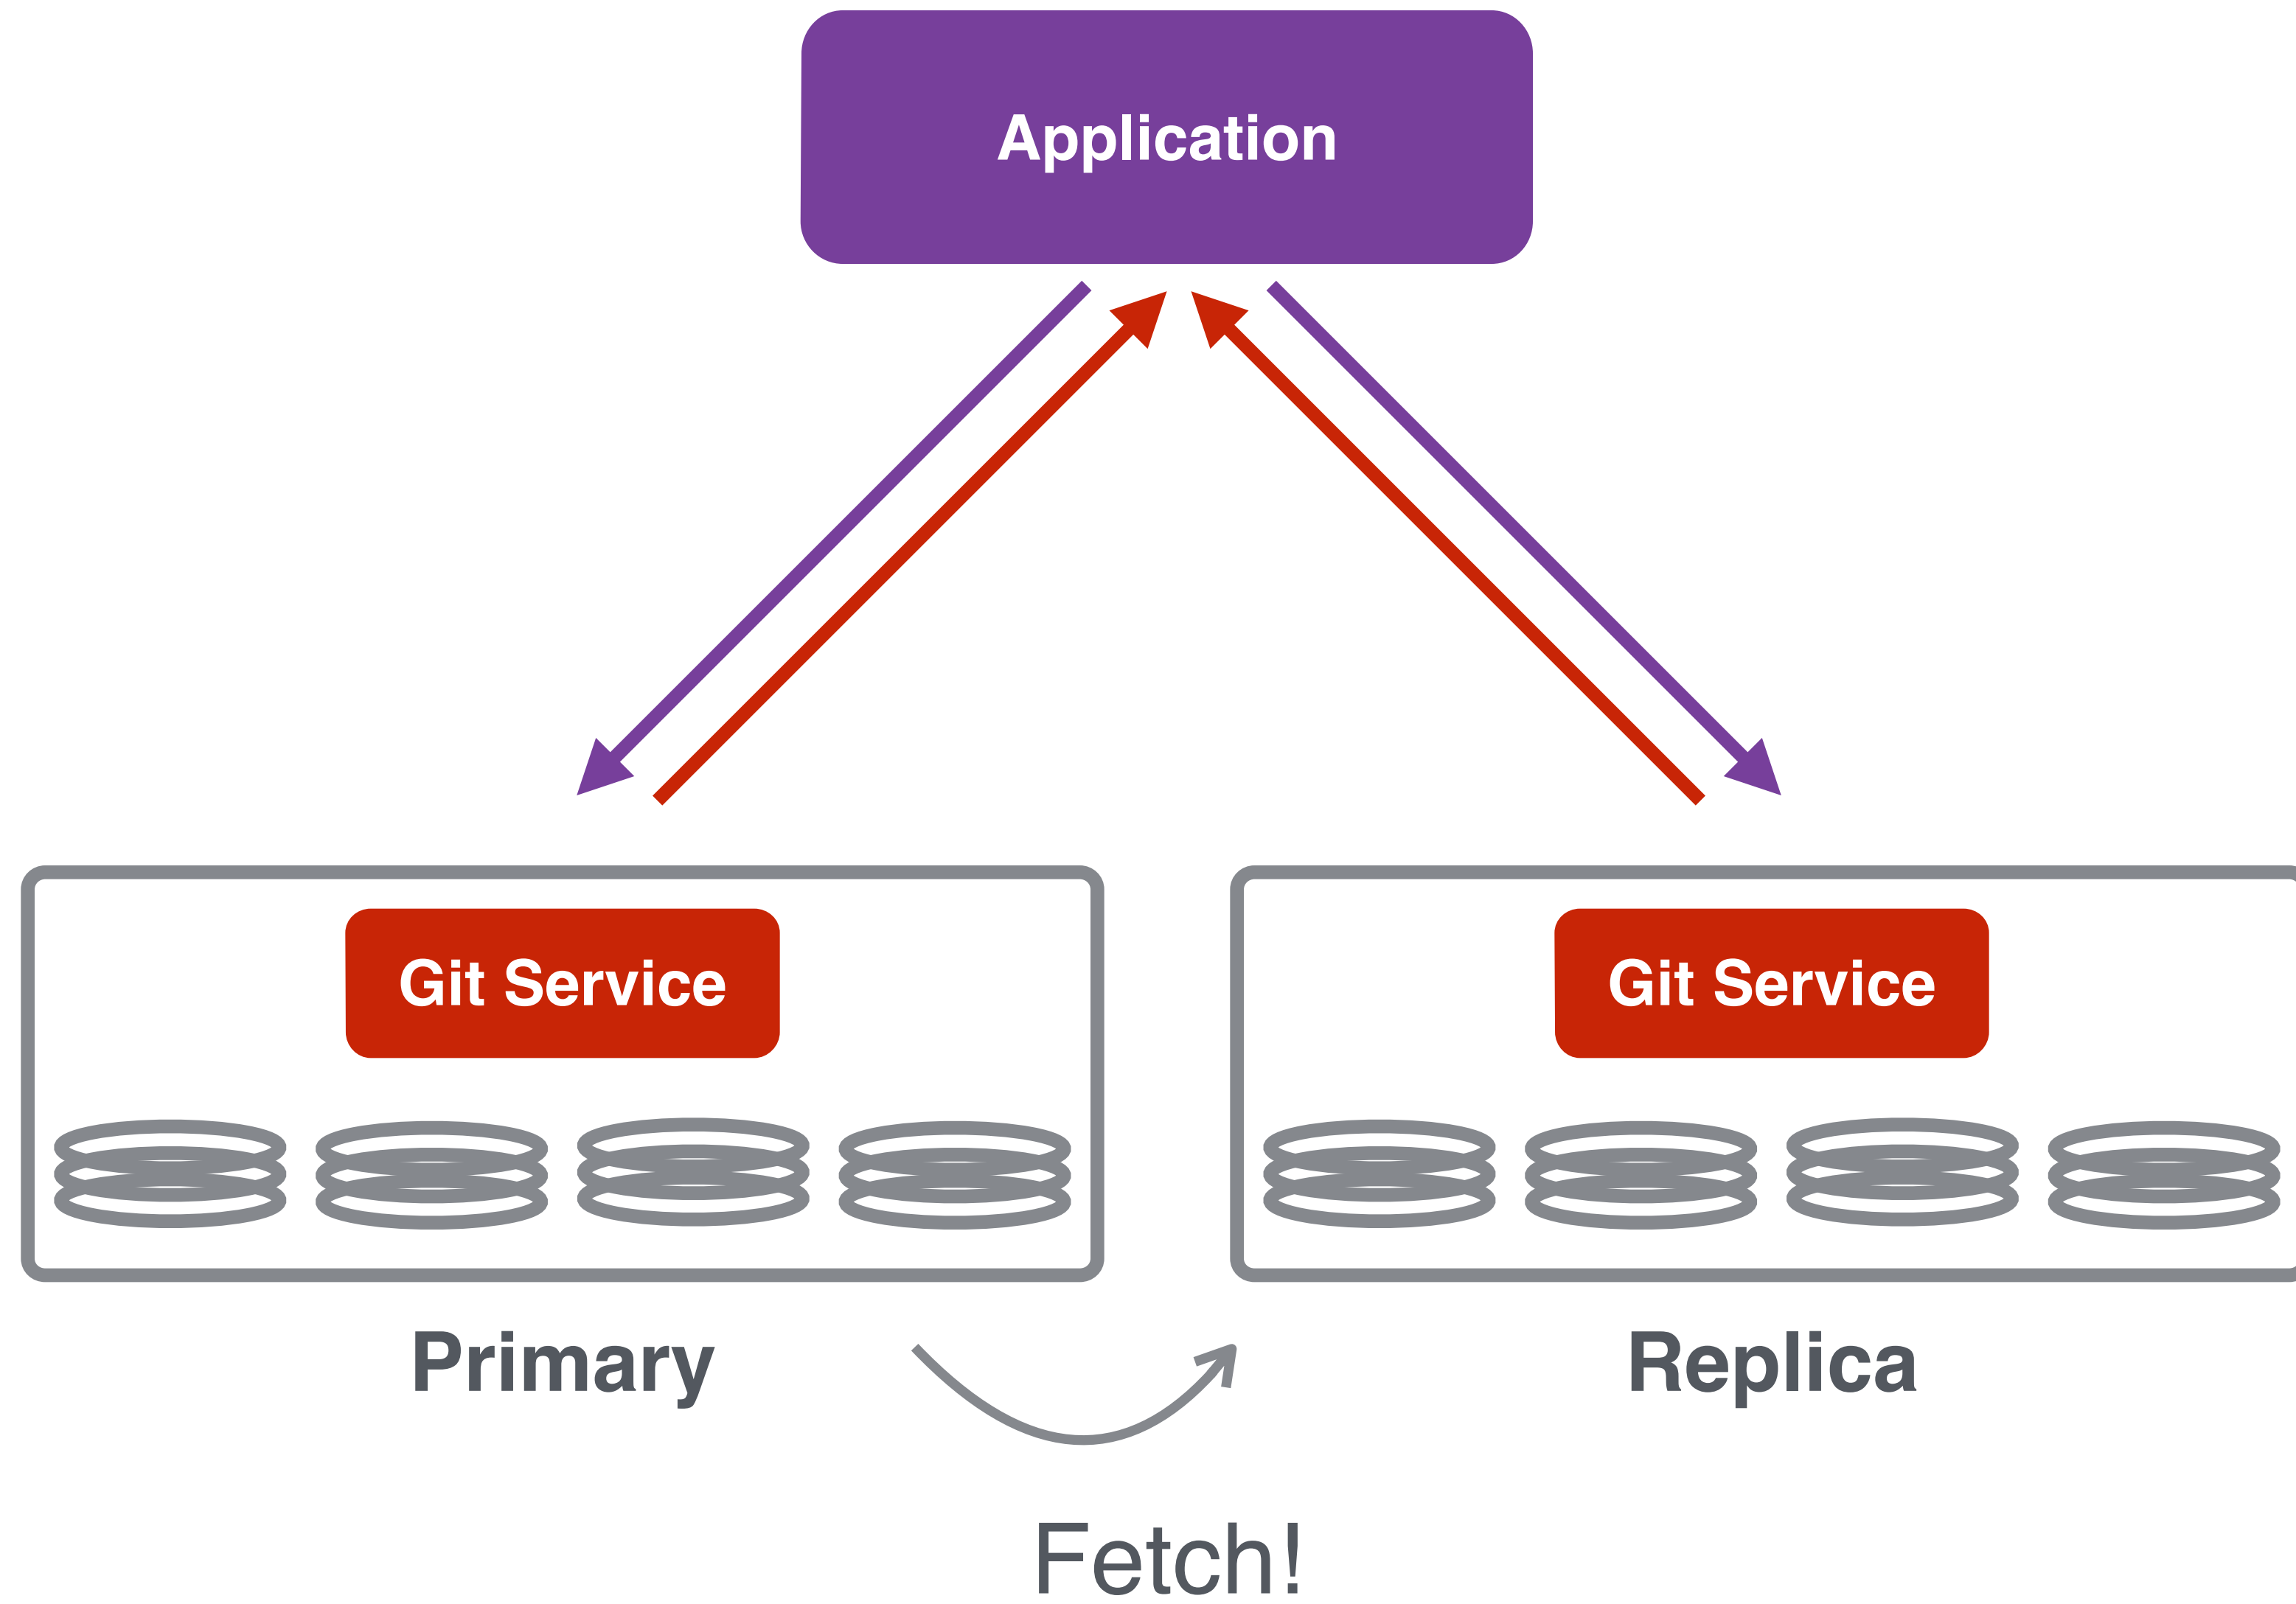
resp, _ := c.Get(url)
```

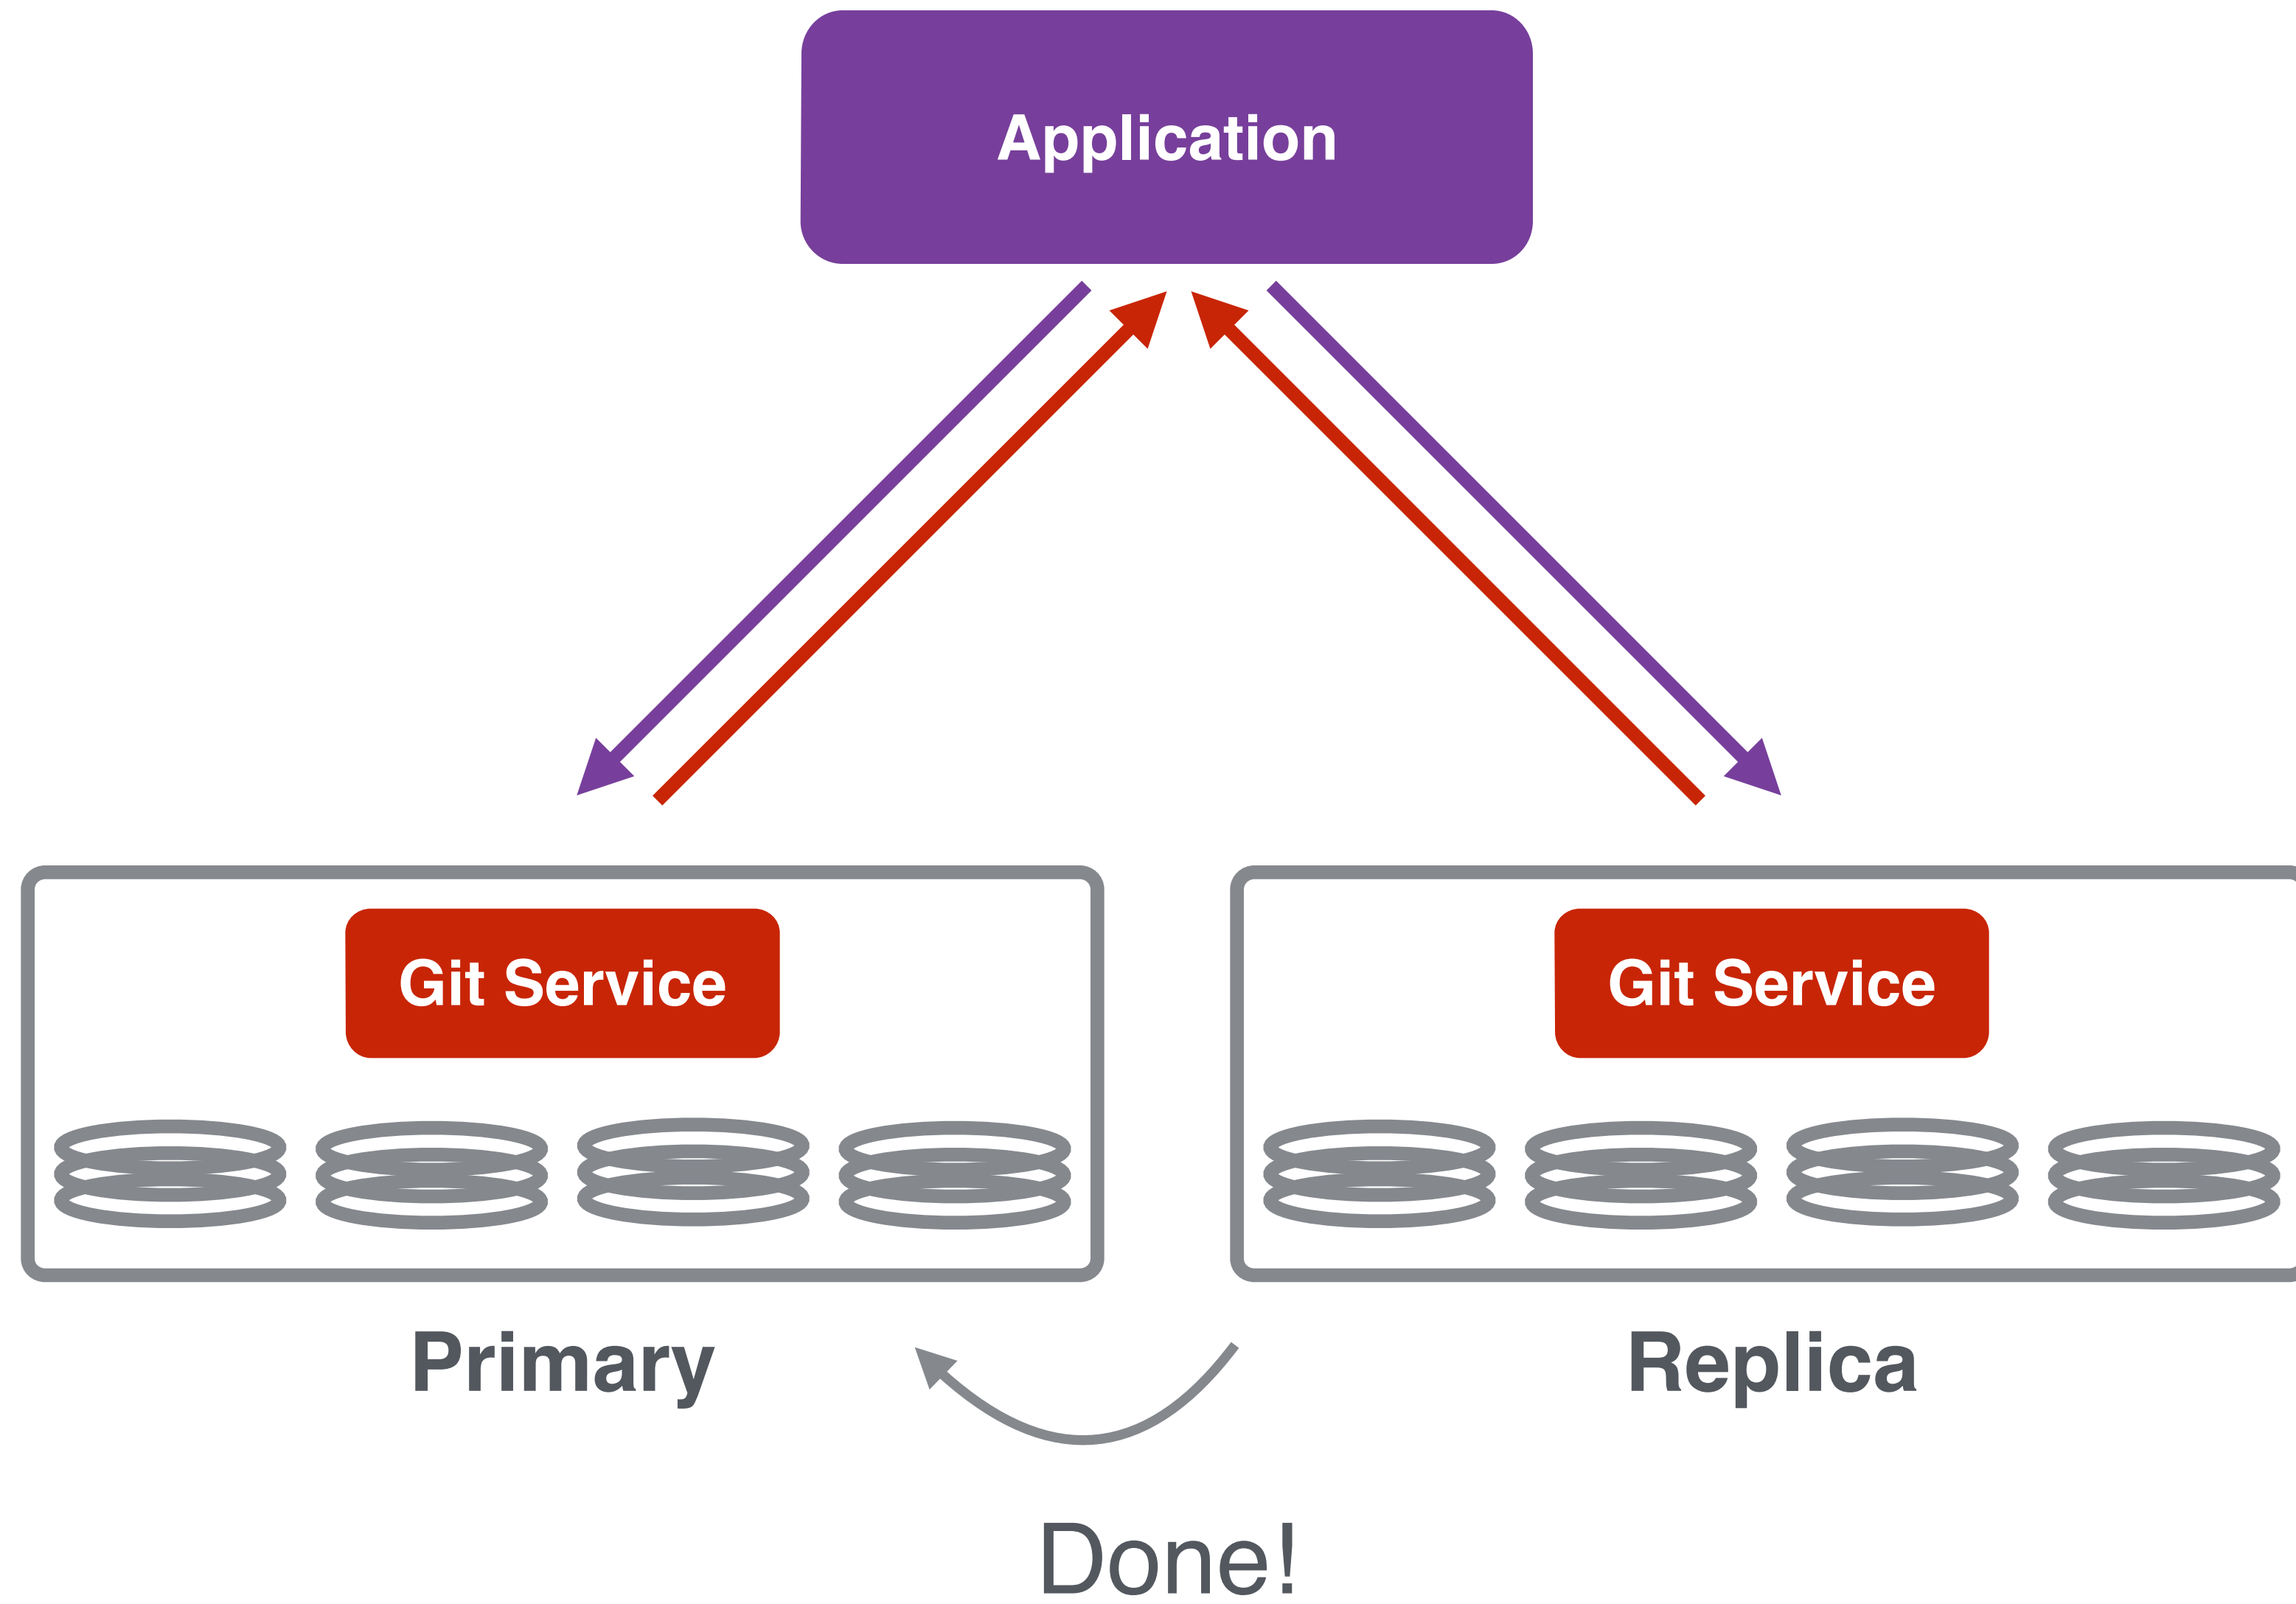
Replication

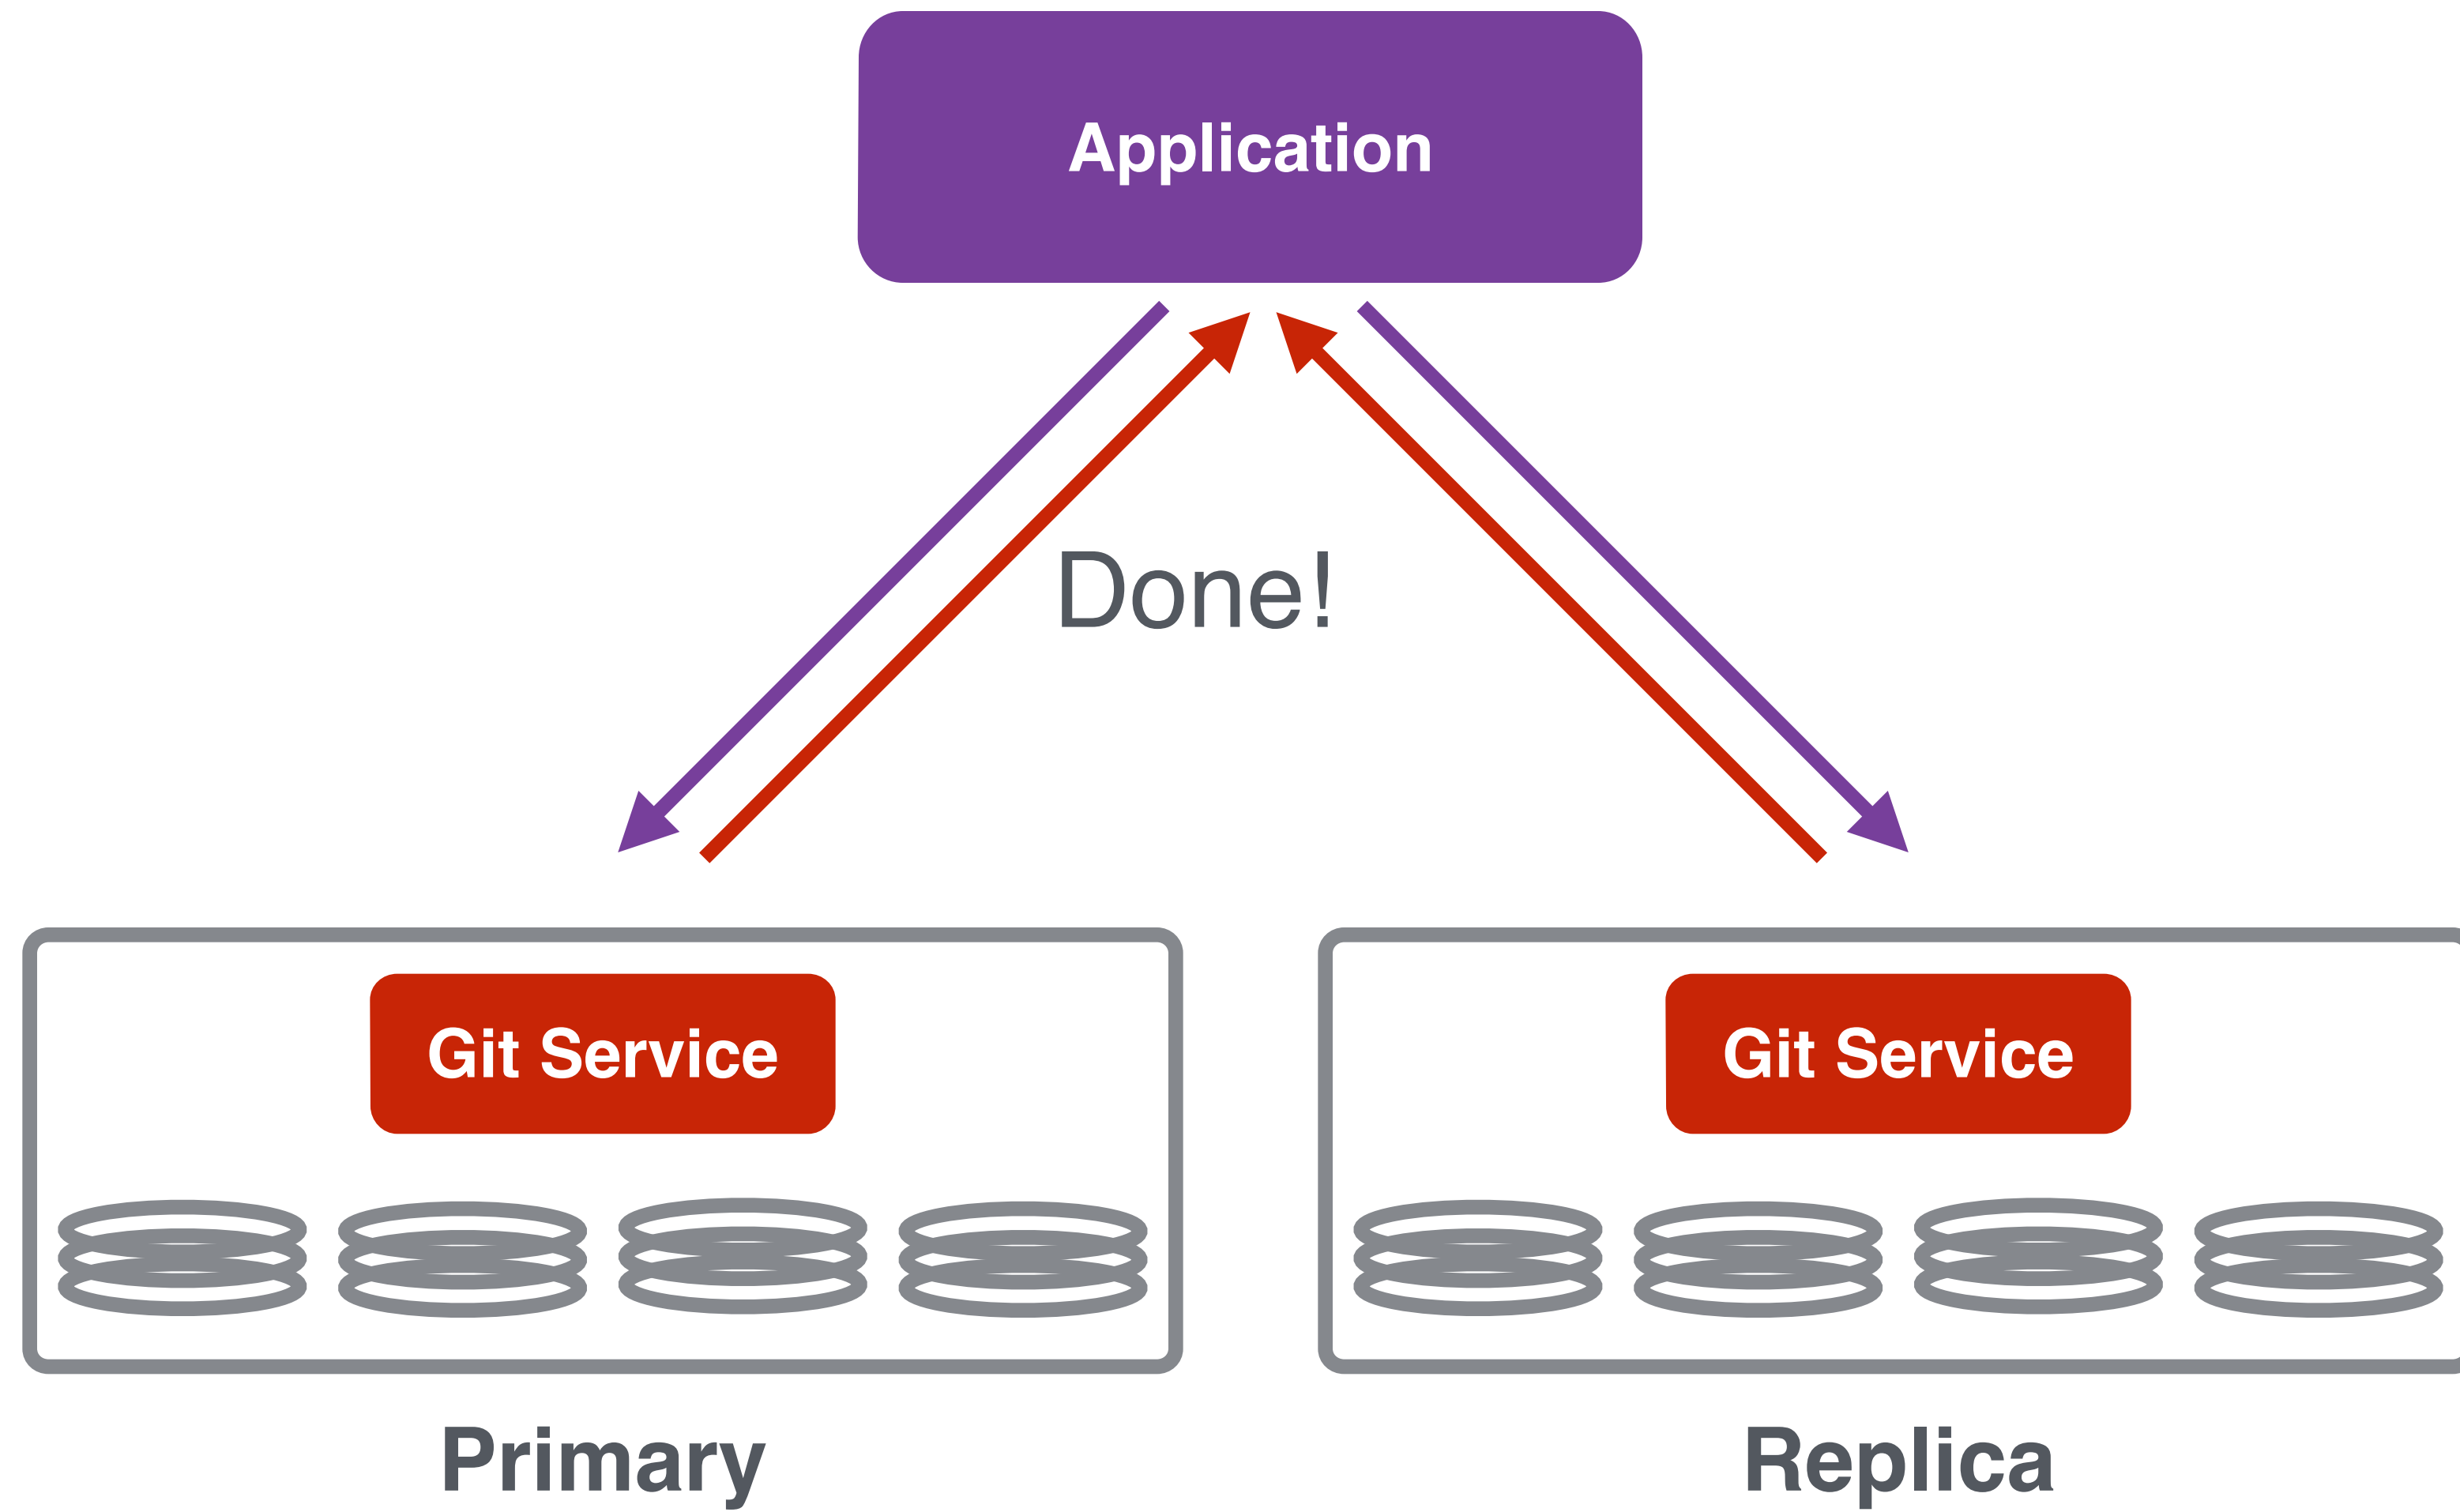
Consistency Vs Availability











```
// Handle fetch requests.  
h := func(w http.ResponseWriter, r *http.Request) {  
    peerChannel := replicateRequest(r)  
  
    rm, _ := repo.LoadRemote("origin")  
    var refsspecs []string  
    rm.Fetch(refspecs, nil, nil)  
  
    if peersChannel != nil {  
        waitForPeers(peerChannel)  
    }  
  
    w.WriteHeader(201)  
}  
  
http.HandleFunc("/r/foo/fetch", h)
```

```
// Replicate request.
func replicateRequest(r *http.Request) chan int {
    if req.Header.Get("X-GIT-REPLICATE") != "" {
        return nil
    }
    peerChannel := make(chan int)

    replicaURL, err := url.Parse(replicaHost)
    replicaURL.Path = r.Path
    replicaURL.Header.Set("X-GIT-REPLICATE", "true")
    req, _ := http.NewRequest("POST", replicaURL.String(), nil)

    go func() {
        resp, _ := httpClient.Do(req)
        peerChannel <- resp.StatusCode
    }()

    return peerChannel
}
```

```
// Wait for replica response.  
func waitForPeers(channel chan int) error {  
    replicaStatus := <- channel  
  
    switch replicaStatus {  
        case 201:  
            // 🎉 🙌 🎉  
        default:  
            // ☔ 😞 ☔  
    }  
  
    return nil  
}
```

git architectures are fun

Thank you!

@calavera