

Concurrent, High-Performance Data-Access With Go

Sunil Sayyaparaju
Engineer
Aerospike

Khosrow Afroozeh
Engineer
Aerospike



What does High Performance mean?

North American RTB speeds & feeds



- 100 millisecond or 150 millisecond ad delivery
 - De-facto standard set in 2004 by Washington Post and others
- North America is 70 to 90 milliseconds wide
 - Two or three data centers
- Newyork-Amsterdam is 150 milliseconds
- Auction is limited to 30 milliseconds
 - Typically closes in 5 milliseconds
- Winners have more data, better models – in 5 milliseconds

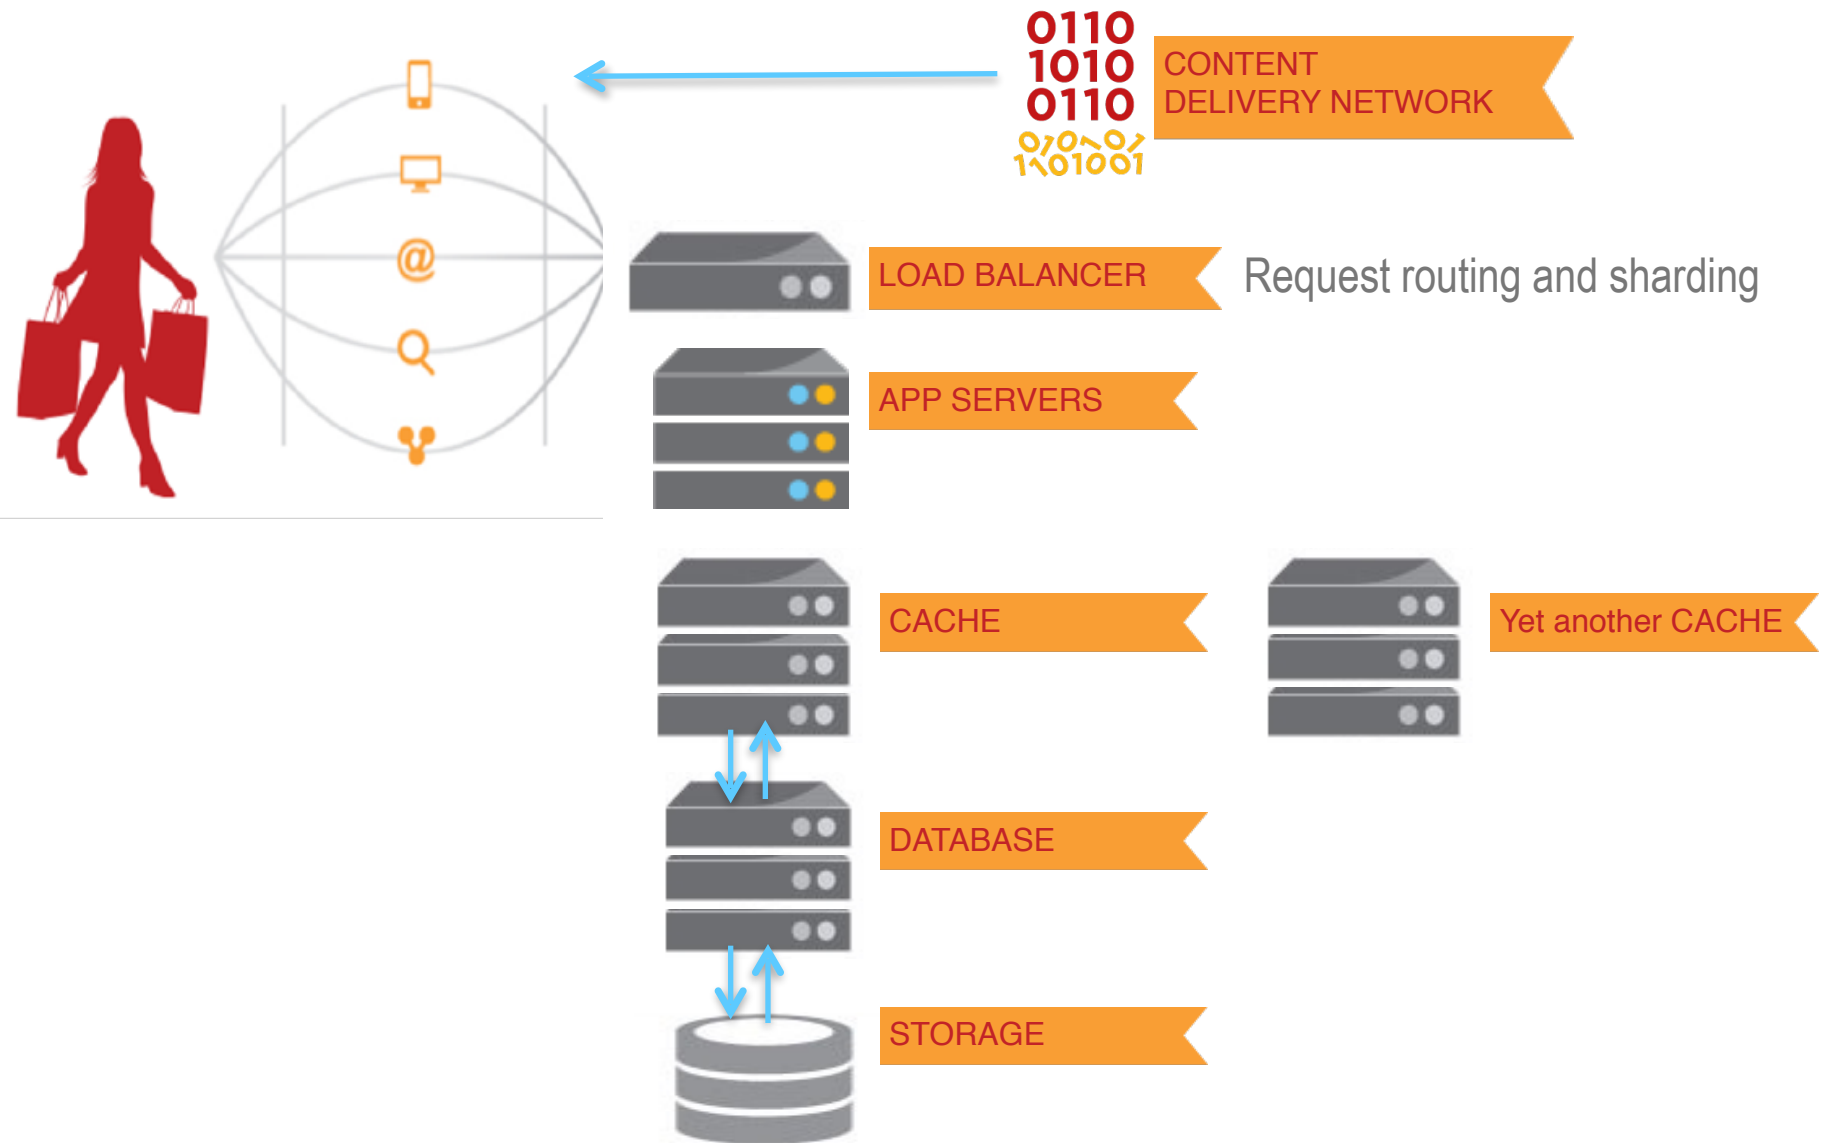
North American RTB speeds & feeds




- 1 to 6 billion cookies tracked
 - Some companies track 200M, some track 20B
- Each bidder has their own data pool
 - Data is your weapon
 - Recent searches, behavior, IP addresses
 - Audience clusters (K-cluster, K-means) from offline Batch Processing
- “Remnant” from Google, Yahoo is about 0.6 million / sec
- Facebook exchange: about 0.6 million / sec
- “others” are 0.5 million / sec

Currently around 2.0M / sec in North America

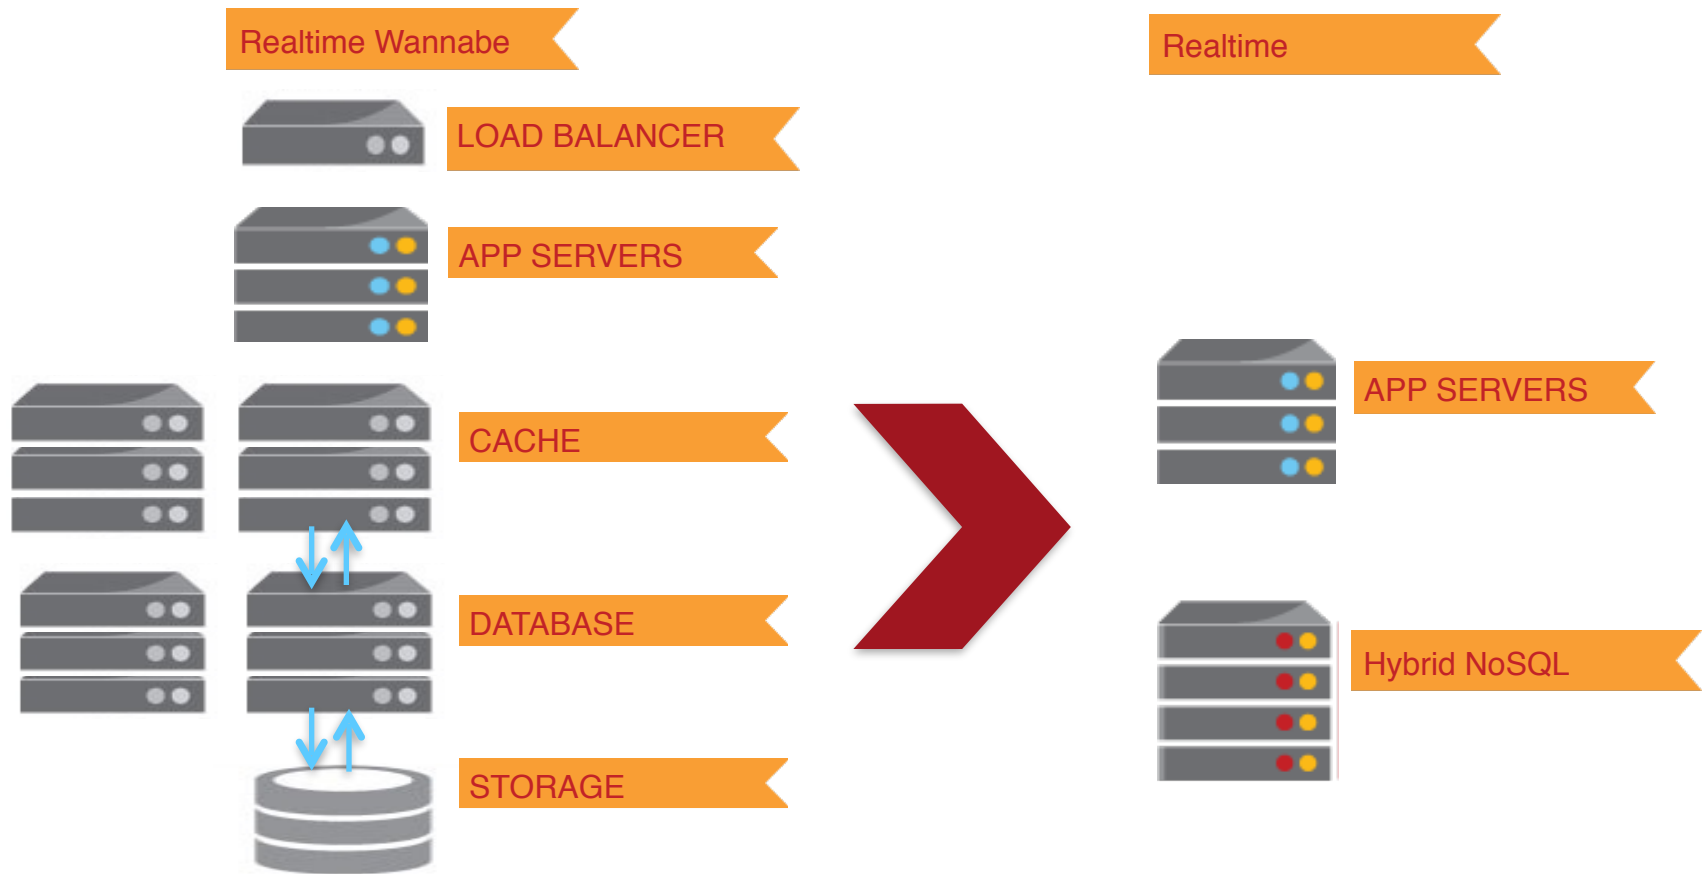
Layer-Driven Architecture



Minimalism Makes a Comeback




Minimalism Makes a Comeback






Aerospike is a distributed database


Hint: It is *NOT* like this!




Hint: It is *NOT* like this!



Aerospike



Aerospike



Benefits of Design:

- 1) **No Hotspots**
– DHT simplifies data partitioning
- 2) **Smart Client** – **1 hop** to data, *no load balancers needed*
- 3) **Shared Nothing** Architecture, every node identical - *no coordinators needed*
- 4) **Single row ACID**
– sync'd replication in cluster
- 5) **Smart Cluster, Zero Touch**
– auto-failover, rebalancing, rack aware, rolling upgrades...
- 6) Transactions and long running tasks prioritized real-time
- 7) **XDR** – sync'd replication across data centers ensures **Zero Downtime**
- 8) **Scale linearly** as data-sizes and workloads increase
- 9) Add capacity with **no service interruption**

Client Libraries Need To Do More - TRANSPARENTLY

- **Implements Aerospike API**

- Optimistic row locking
- Optimized binary protocol

- **Cluster Tracking**

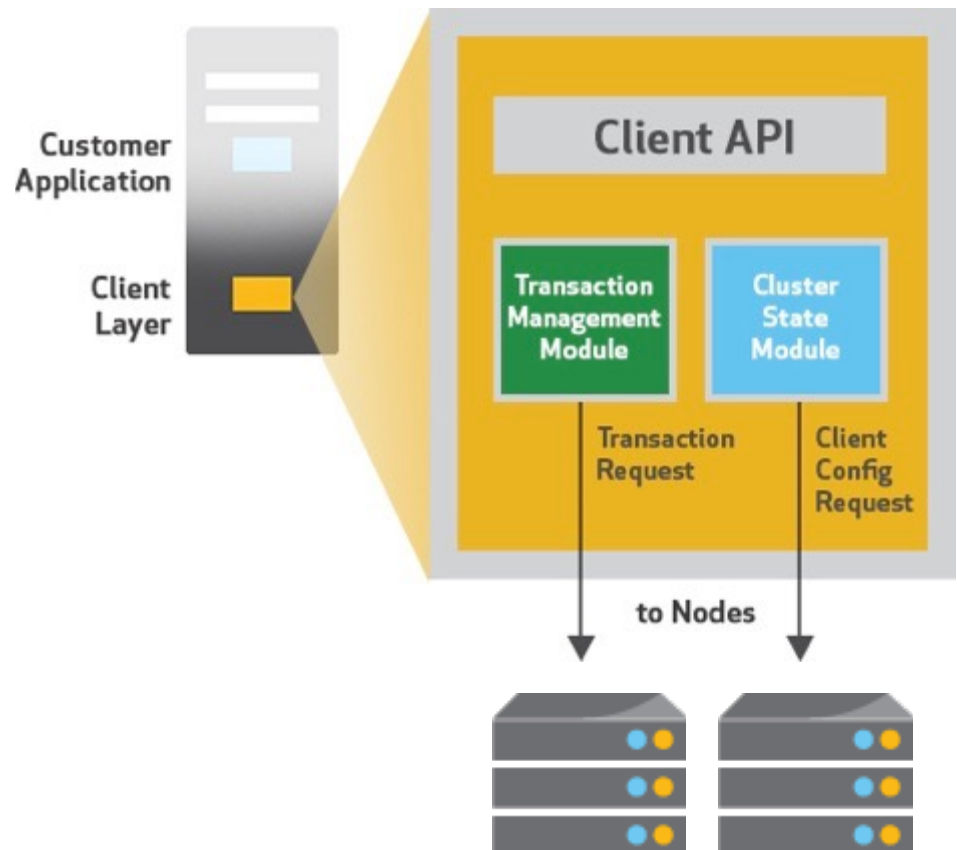
- Learns about cluster changes, Partition Map

- **Transaction Semantics**

- Global Transaction ID
- Retransmit and timeout

- **Linearly Scales**


- No extra hop to data
- No load balancers in the way





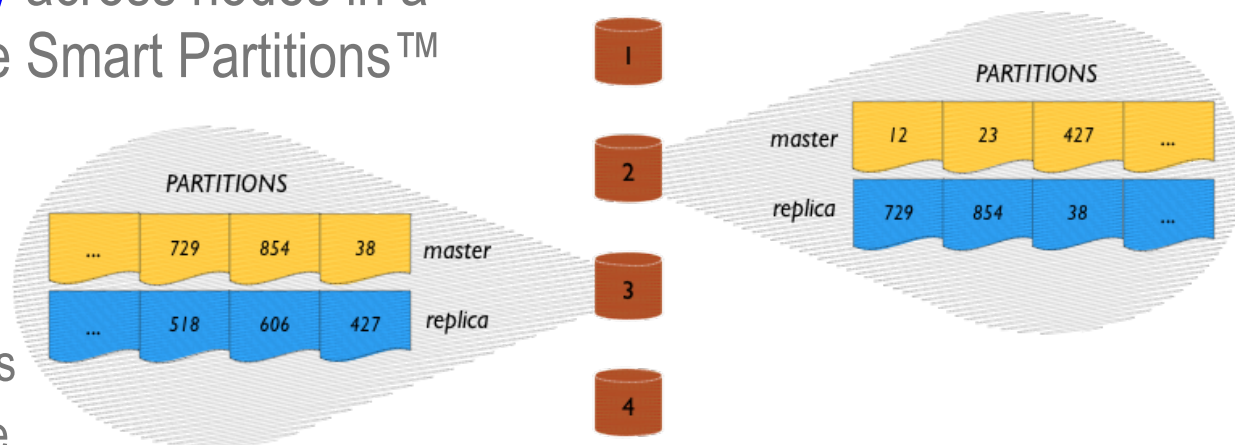
Tracking Cluster Changes

Cluster Tracking




Data is **distributed evenly** across nodes in a cluster using the Aerospike Smart Partitions™ algorithm.

- Even distribution of
 - **Partitions** across nodes
 - **Records** across Partitions
 - **Data** across Flash device



Cluster Tracking

Spawn Goroutine To
Track Cluster Changes



```
func NewCluster(policy *ClientPolicy, hosts []*Host) (*Cluster, error) {
    newCluster := &Cluster{...}

    // start up cluster maintenance go routine
    newCluster.wgTend.Add(1)
    go newCluster.clusterBoss(policy)

    return newCluster, nil
}

func (clstr *Cluster) clusterBoss(policy *ClientPolicy) {
    defer clstr.wgTend.Done()

Loop:
    for {
        select {
        case <-clstr.tendChannel:
            // tend channel closed
            break Loop
        case <-time.After(tendInterval):
            if err := clstr.tend(); err != nil {
                Logger.Warn(err.Error())
            }
        }
    }

    // cleanup code goes here
    // close the nodes, ...
}

func (clstr *Cluster) Close() {
    if !clstr.closed.Get() {
        // send close signal to maintenance channel
        close(clstr.tendChannel)

        // wait until tend is over
        clstr.wgTend.Wait()
    }
}
```


Cluster Tracking

On Intervals, Update
Cluster Status

```
func NewCluster(policy *ClientPolicy, hosts []*Host) (*Cluster, error) {
    newCluster := &Cluster{...}

    // start up cluster maintenance go routine
    newCluster.wgTend.Add(1)
    go newCluster.clusterBoss(policy)

    return newCluster, nil
}

func (clstr *Cluster) clusterBoss(policy *ClientPolicy) {
    defer clstr.wgTend.Done()

Loop:
    for {
        select {
        case <-clstr.tendChannel:
            // tend channel closed
            break Loop
        case <-time.After(tendInterval):
            if err := clstr.tend(); err != nil {
                Logger.Warn(err.Error())
            }
        }
    }

    // cleanup code goes here
    // close the nodes, ...
}

func (clstr *Cluster) Close() {
    if !clstr.closed.Get() {
        // send close signal to maintenance channel
        close(clstr.tendChannel)

        // wait until tend is over
        clstr.wgTend.Wait()
    }
}
```

Cluster Tracking

```
func NewCluster(policy *ClientPolicy, hosts []*Host) (*Cluster, error) {
    newCluster := &Cluster{...}

    // start up cluster maintenance go routine
    newCluster.wgTend.Add(1)
    go newCluster.clusterBoss(policy)

    return newCluster, nil
}

func (clstr *Cluster) clusterBoss(policy *ClientPolicy) {
    defer clstr.wgTend.Done()


Loop:
    for {
        select {
        case <-clstr.tendChannel:
            // tend channel closed
            break Loop
        case <-time.After(tendInterval):
            if err := clstr.tend(); err != nil {
                Logger.Warn(err.Error())
            }
        }
    }

    // cleanup code goes here
    // close the nodes, ...
}

func (clstr *Cluster) Close() {
    if !clstr.closed.Get() {
        // send close signal to maintenance channel
        close(clstr.tendChannel)

        // wait until tend is over
        clstr.wgTend.Wait()
    }
}
```

Broadcast Closing Of
Cluster



Cluster Tracking

Break the loop to clean up

```
func NewCluster(policy *ClientPolicy, hosts []*Host) (*Cluster, error) {
    newCluster := &Cluster{...}

    // start up cluster maintenance go routine
    newCluster.wgTend.Add(1)
    go newCluster.clusterBoss(policy)

    return newCluster, nil
}

func (clstr *Cluster) clusterBoss(policy *ClientPolicy) {
    defer clstr.wgTend.Done()

Loop:
    for {
        select {
        case <-clstr.tendChannel:
            // tend channel closed
            break Loop
        case <-time.After(tendInterval):
            if err := clstr.tend(); err != nil {
                Logger.Warn(err.Error())
            }
        }
    }

    // cleanup code goes here
    // close the nodes, ...
}


func (clstr *Cluster) Close() {
    if !clstr.closed.Get() {
        // send close signal to maintenance channel
        close(clstr.tendChannel)

        // wait until tend is over
        clstr.wgTend.Wait()
    }
}
```



How To Stream Back Data From Multiple Nodes

Queries



**Queries are Scatter
And Gather**

1 Goroutine Per Node

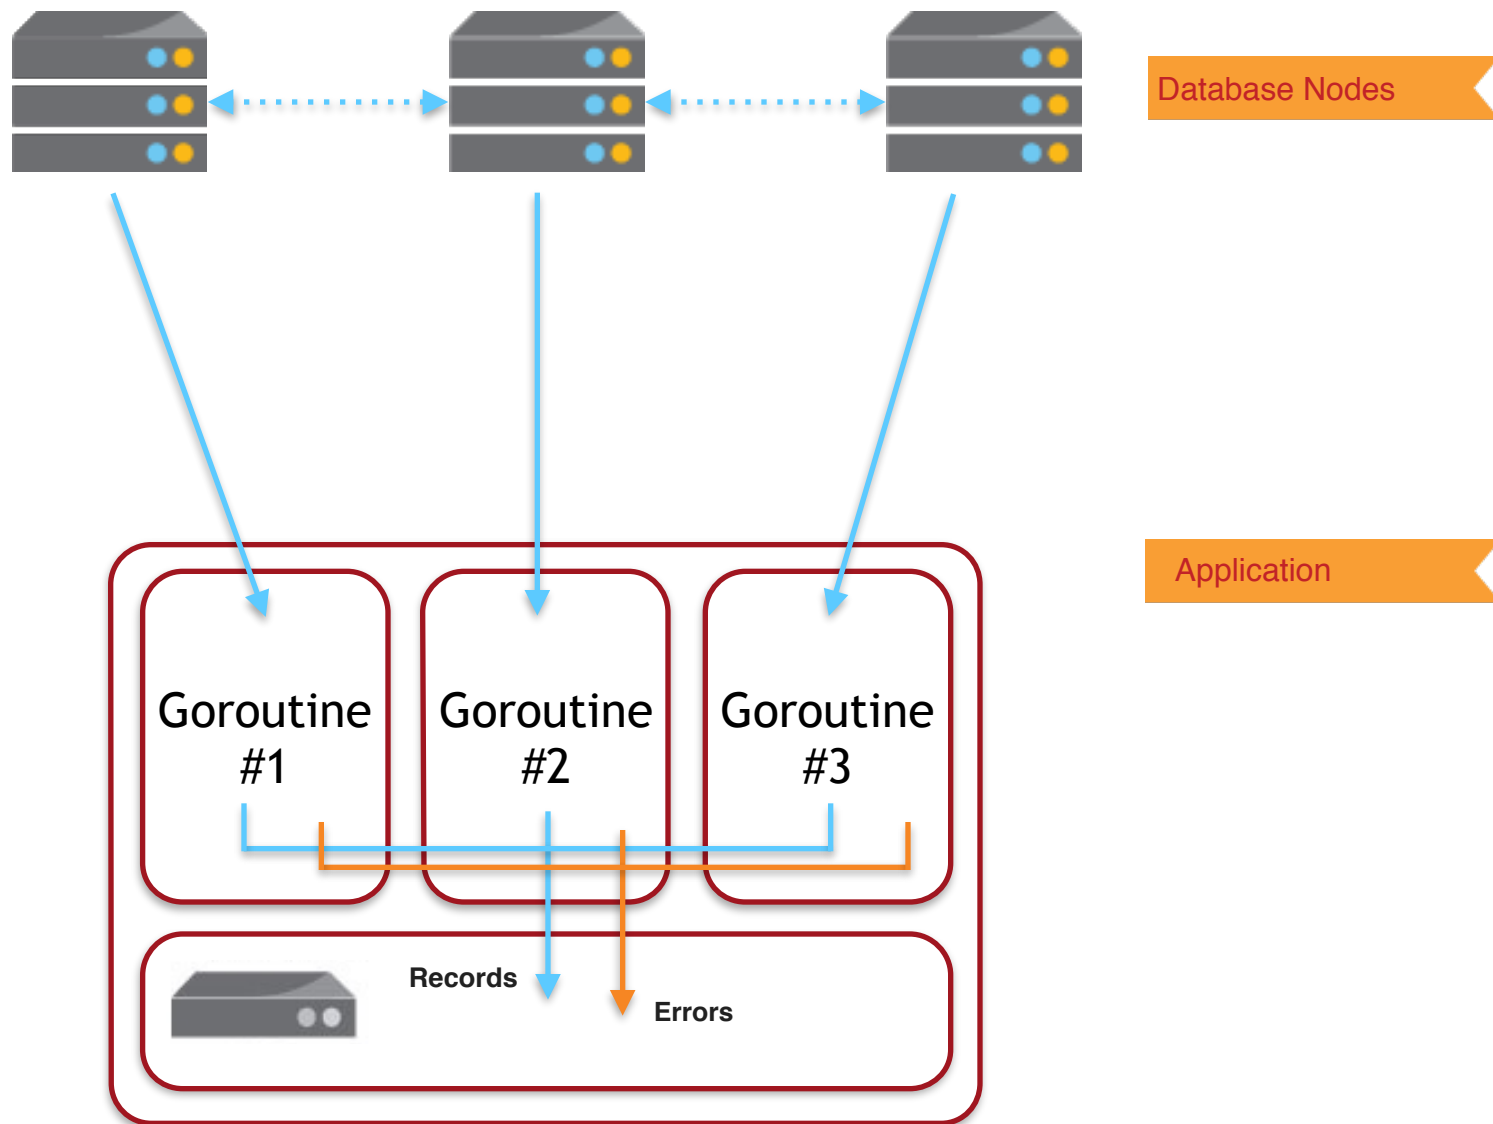
**Data Is Streamed
Back**

No Cursors

Errors Are Per Node

Might Want To
Continue With Other
Nodes

Queries



Queries

```
// Recordset encapsulates the result of Scan and Query commands.
type Recordset struct {
    Records chan *Record
    Errors  chan error

    wgGoroutines sync.WaitGroup
    goroutines   *AtomicInt

    active    *AtomicBool
    cancelled chan struct{}
}

// Close all streams from different nodes.
func (rcs *Recordset) Close() {
    // No panic on several calls
    if rcs.active.CompareAndToggle(true) {
        // broadcast to all command goroutines listening to the channel
        close(rcs.cancelled)

        // wait until all goroutines are done
        rcs.wgGoroutines.Wait()

        close(rcs.Records)
        close(rcs.Errors)
    }
}

func (rcs *Recordset) signalEnd() {
    rcs.wgGoroutines.Done()
    if rcs.goroutines.DecrementAndGet() == 0 {
        rcs.Close()
    }
}
```

```
func (cmd *scanCommand) parseRecordResults(ifc command, receiveSize int) (bool, error) {
    ...
    for cmd.dataOffset < receiveSize {
        ...
        if err := cmd.readBytes(int(_MSG_REMAINING_HEADER_SIZE)); err != nil {
            cmd.recordset.Errors <- newNodeError(cmd.node, err)
            return false, err
        }

        // parse record

        // If the channel is full and it blocks, we don't want this command to
        // block forever
        select {
        // send back the result on the buffered channel
        case cmd.recordset.Records <- newRecord(cmd.node, key, bins, generation, expiration):
        case <-cmd.recordset.cancelled:
            return false, NewAerospikeError(SCAN_TERMINATED)
        }

        // all records streamed successfully
        return true, nil
    }
}

func (cmd *scanCommand) Execute() error {
    defer cmd.recordset.signalEnd()
    ...
}
```

Same buffered channel is used for all goroutines

Queries

```
// Recordset encapsulates the result of Scan and Query commands.
type Recordset struct {
    Records chan *Record
    Errors chan error
    wgGoroutines sync.WaitGroup
    goroutines *AtomicInt

    active *AtomicBool
    cancelled chan struct{}
}

// Close all streams from different nodes.
func (rcs *Recordset) Close() {
    // No panic on several calls
    if rcs.active.CompareAndToggle(true) {
        // broadcast to all command goroutines listening to the channel
        close(rcs.cancelled)

        // wait until all goroutines are done
        rcs.wgGoroutines.Wait()

        close(rcs.Records)
        close(rcs.Errors)
    }
}

func (rcs *Recordset) signalEnd() {
    rcs.wgGoroutines.Done()
    if rcs.goroutines.DecrementAndGet() == 0 {
        rcs.Close()
    }
}


func (cmd *scanCommand) parseRecordResults(ifc command, receiveSize int) (bool, error) {
    ...
    for cmd.dataOffset < receiveSize {
        ...
        if err := cmd.readBytes(int(_MSG_REMAINING_HEADER_SIZE)); err != nil {
            cmd.recordset.Errors <- newNodeError(cmd.node, err)
            return false, err
        }

        // parse record

        // If the channel is full and it blocks, we don't want this command to
        // block forever
        select {
        // send back the result on the buffered channel
        case cmd.recordset.Records <- newRecord(cmd.node, key, bins, generation, expiration):
        case <-cmd.recordset.cancelled:
            return false, NewAerospikeError(SCAN_TERMINATED)
        }

        // all records streamed successfully
        return true, nil
    }
}

func (cmd *scanCommand) Execute() error {
    defer cmd.recordset.signalEnd()
    ...
}
```



Errors are sent back on a separate channel


```
// Recordset encapsulates the result of Scan and Query commands.
type Recordset struct {
    Records chan *Record
    Errors  chan error

    wgGoroutines sync.WaitGroup
    goroutines   *AtomicInt

    active    *AtomicBool
    cancelled chan struct{}
}

// Close all streams from different nodes.
func (rcs *Recordset) Close() {
    // No panic on several calls
    if rcs.active.CompareAndToggle(true) {
        // broadcast to all command goroutines listening to the channel
        close(rcs.cancelled)

        // wait until all goroutines are done
        rcs.wgGoroutines.Wait()

        close(rcs.Records)
        close(rcs.Errors)
    }
}

func (rcs *Recordset) signalEnd() {
    rcs.wgGoroutines.Done()
    if rcs.goroutines.DecrementAndGet() == 0 {
        rcs.Close()
    }
}
```

```
func (cmd *scanCommand) parseRecordResults(ifc command, receiveSize int) (bool, error) {
    ...
    for cmd.dataOffset < receiveSize {
        ...
        if err := cmd.readBytes(int(_MSG_REMAINING_HEADER_SIZE)); err != nil {
            cmd.recordset.Errors <- newNodeError(cmd.node, err)
            return false, err
        }

        // parse record

        // If the channel is full and it blocks, we don't want this command to
        // block forever
        select {
            // send back the result on the buffered channel
            case cmd.recordset.Records <- newRecord(cmd.node, key, bins, generation, expiration):
            case <-cmd.recordset.cancelled:
                return false, NewAerospikeError(SCAN_TERMINATED)
        }

        // all records streamed successfully
        return true, nil
    }
}

func (cmd *scanCommand) Execute() error {
    defer cmd.recordset.signalEnd()
    ...
}
```

Defer on all producers will ensure **Close()** is called at least once

Queries

Will not deadlock, all producers will return

```
// Close all streams from different nodes.
func (rcs *Recordset) Close() {
    // No panic on several calls
    if rcs.active.CompareAndToggle(true) {
        // broadcast to all command goroutines listening to the channel
        close(rcs.cancelled)

        // wait until all goroutines are done
        rcs.wgGoroutines.Wait()

        close(rcs.Records)
        close(rcs.Errors)
    }
}

func (rcs *Recordset) signalEnd() {
    rcs.wgGoroutines.Done()
    if rcs.goroutines.DecrementAndGet() == 0 {
        rcs.Close()
    }
}
```

Consumer

```
// If the channel is full and it blocks, we don't want this command to
// block forever
select {
    // send back the result on the buffered channel
    case cmd.recordset.Records <- newRecord(cmd.node, key, bins, generation, expiration):
    case <-cmd.recordset.cancelled:
        return false, NewAerospikeError(SCAN_TERMINATED)
}
```

Producer #N

The Above code will never panic, or block indefinitely:

- **defer** on **Execute()** will ensure **Close()** is called at least once
- When **cancelled** channel is closed, even if the **Records** channel is full and blocked, **select** will go through and return from goroutine

Queries

Will not panic, all producers have already stopped

```
// Close all streams from different nodes.
func (rcs *Recordset) Close() {
    // No panic on several calls
    if rcs.active.CompareAndToggle(true) {
        // broadcast to all command goroutines listening to the channel
        close(rcs.cancelled)

        // wait until all goroutines are done
        rcs.wgGoroutines.Wait()

        {
            close(rcs.Records)
            close(rcs.Errors)
        }
    }

    func (rcs *Recordset) signalEnd() {
        rcs.wgGoroutines.Done()
        if rcs.goroutines.DecrementAndGet() == 0 {
            rcs.Close()
        }
    }
}
```

Consumer

```
// If the channel is full and it blocks, we don't want this command to
// block forever
select {
    // send back the result on the buffered channel
    case cmd.recordset.Records <- newRecord(cmd.node, key, bins, generation, expiration):
    case <-cmd.recordset.cancelled:
        return false, NewAerospikeError(SCAN_TERMINATED)
}
```

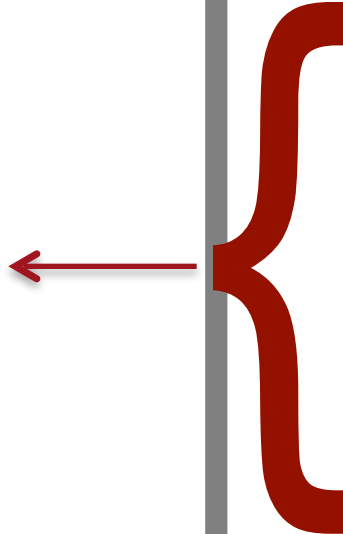
Producer #N

The Above code will never panic, or block indefinitely:

- **defer** on **Execute()** will ensure **Close()** is called at least once
- When **cancelled** channel is closed, even if the **Records** channel is full and blocked, **select** will go through and return from goroutine

Consumption Pattern

Supports
numerous
consumers




```
func (rcs *Recordset) Results() <-chan *result {
    res := make(chan *result, len(rcs.Records))

    go func() {
        L:
        for {
            select {
            case r := <-rcs.Records:
                if r != nil {
                    res <- &result{Record: r, Err: nil}
                } else {
                    close(res)
                    break L
                }
            case e := <-rcs.Errors:
                if e != nil {
                    res <- &result{Record: nil, Err: e}
                }
            }
        }
    }()

    return (<-chan *result)(res)
}
```

Fork #N

Range on
Channel, get
either a **Record**
or **error**



```
recordset, err := client.ScanAll(...)
for res := range recordset.Results() {
    if res.Err != nil {
        // handle error here
    } else {
        // process record here
    }
}
```

Consumer

Errors do not mean the stream has ended. They carry information about *what* went wrong, and *where*

Take away

- **Channels are a good abstraction for data streaming**
 - Easier to avoid deadlocks
- **Use the least number of channels and goroutines**
 - Multiplex on demand
 - Goroutines are cheap, but not free
 - Pass channels throughout API, don't ask for new ones



How To Win Against The Garbage Collector

Don't Create Garbage

- **Avoid Memory Allocations**

- Might be trickier than you think!

- **Pool All Objects**

- Buffers
- Hash Objects
- Network Connections
- Reusable Intermediate / Temporary Objects
- Random Number Generators

- `sync.Pool` **is not ideal to pool long living objects**

- Integrated into GC, sweeps and deallocates objects too often

- **LIFO pools tend to do better regarding CPU cache**

- Channels for pool implementation still perform well though

Don't Reallocate Memory

- **Avoid Memory Allocations**
 - Might be trickier than you think!
- **All following code snippets allocate memory:**

```
b := new(bytes.Buffer)
longBuf := bytes.Repeat([]byte{32}, 100)

// Allocation happens in Write to Grow the buffer
b.Write([]byte{32})

// Re-allocation happens in Write to Grow the buffer
b.Write(longBuf)
```

```
// allocation happens inside Trim
strings.Trim(" Trim this string ", " ")
```

```
h := ripemd160.New()
h.Write(data)

// allocation happens inside Sum()
h.Sum(nil)

// Why? To support the general case
func (d *digest) Sum(in []byte) []byte {
    // Make a copy of d0 so that caller can keep writing and summing.
    d := *d0
```


Don't Reallocate Memory



bradfitz commented on dd95d39 on Dec 9, 2014



Why use `bytes.NewBuffer` at all? The zero value is smaller and usable and more idiomatic.



khaf commented on dd95d39 on Dec 9, 2014

Collaborator



The answer is:

1. Excessive re-allocation during growth
2. Usage patterns of people using Key/Value stores.

Complex objects (Maps and Lists) are used excessively, and the buffer invariably ends up growing (re-allocating) more than a few times if it is too small.

There's no magic number that is good enough, and I'm not a fan of heuristics in libraries, so I'll probably end up redesigning to pool packers somehow.

Thoughts?

Indeterministic Pools are Resource Leaks You Haven't Found Yet! *

* Please apply wisdom to the quip above

Pools

- Make Pools Deterministic
- Let User Decide About Parameters Per Use-Case
- Document The Above!

Set & Enforce Limits for
Pooled Buffers

```
type BufferPool struct {  
    pool    [][]byte  
    poolSize int  
  
    pos int64  
  
    {  
        maxBufSize int  
        initBufSize int  
    }  
  
    mutex sync.Mutex  
}  
  
func (bp *BufferPool) Put(buf []byte) {  
    {  
        if len(buf) <= bp.maxBufSize {  
            ...  
        }  
    }  
    // buffer too big to go back into the pool  
}
```

Set Sensible Defaults

```
// a custom buffer pool with fine grained control over its contents  
// maxSize: 128KiB  
// initial bufferSize: 16 KiB  
// maximum buffer size to keep in the pool: 128K  
var bufPool = NewBufferPool(512, 16*1024, 128*1024)  
  
// SetCommandBufferPool can be used to customize the command Buffer Pool parameter  
// the pool for different workloads  
func SetCommandBufferPool(poolSize, initBufSize, maxBufferSize int) {  
    bufPool = NewBufferPool(poolSize, initBufSize, maxBufferSize)  
}
```

Let Users Tweak Per
Use-Case

Non-Blocking Sync'd Queues Using Channels

Non-blocking Channel Read / Write Is Possible

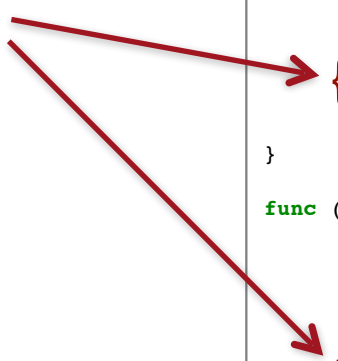
Non-Blocking

```
type AtomicQueue struct {
    items chan interface{}
}

func NewAtomicQueue(size int) *AtomicQueue {
    return &AtomicQueue{
        items: make(chan interface{}, size),
    }
}

func (aq *AtomicQueue) Offer(item interface{}) bool {
    // non-blocking send pattern
    select {
    case aq.items <- item:
        return true
    default:
    }
    return false
}

func (aq *AtomicQueue) Poll() interface{} {
    // non-blocking read pattern
    select {
    case item := <-aq.items:
        return item
    default:
    }
    return nil
}
```



Generic Pool Using Sync'd Queue

```
type Pool struct {
    pool *AtomicQueue

    // New will create a new object if pool is empty
    New func(params ...interface{}) interface{}

    // StillUsable checks if the object polled from the pool
    // is still fresh and usable
    StillUsable func(obj interface{}, params ...interface{}) bool

    // CanReturn checks if the object is eligible to go back to the pool
    CanReturn func(obj interface{}) bool

    // Finalize will be called when an object is not
    // eligible to go back to the pool.
    // Usable to close connections, file handles, ...
    Finalize func(obj interface{})
}
```



How Does It Perform?

Comparison vs C and Java

Data Type: String of size 2000

Transactions Per Second

	<i>100% WRITE</i>	<i>100% READ</i>	<i>50% READ 50% WRITE</i>
<i>C</i>	261K	308K	273K
<i>JAVA</i>	180K	280K	220K
<i>GO</i>	306K	238K	276K

 Fastest

 Slowest

Comparison vs C and Java


Data Type: Integer (8 Bytes)

Transactions Per Second

	<i>100% WRITE</i>	<i>100% READ</i>	<i>50% READ 50% WRITE</i>
<i>C</i>	<i>299K</i>	<i>334K</i>	<i>307K</i>
<i>JAVA</i>	<i>308K</i>	<i>339K</i>	<i>325K</i>
<i>GO</i>	<i>338K</i>	<i>328K</i>	<i>330K</i>

 Fastest

 Slowest



Thank you

@sunilvirus

sunil@aerospike.com

@parshua

khosrow@aerospike.com